



Numerical linear algebra in the Canadian hydrological model

Kevin R. Green

Department of Computer Science
University of Saskatchewan

Math 313
University of Saskatchewan
September 15, 2020

Outline

1 Hydrology

- Overview
- CHM
- Physical processes

2 Space-dependent processes

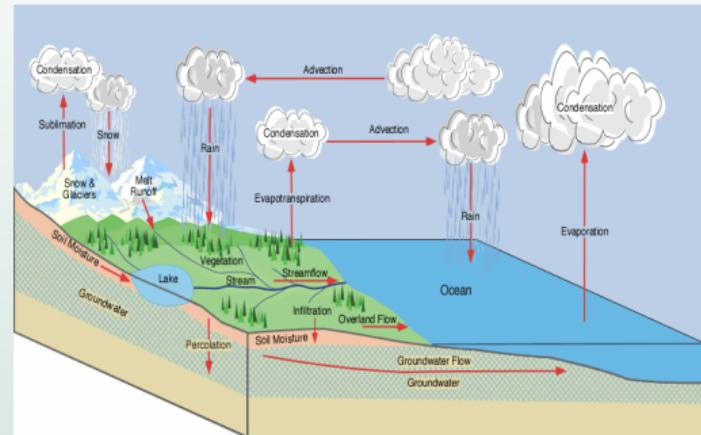
- Partial differential equations
- Mesh triangulation
- Discretization

3 Linear algebra

- Solvers
- Permutations
- Experiments

Hydrology

- The scientific study of the movement, distribution, and quality of water on Earth and other planets.

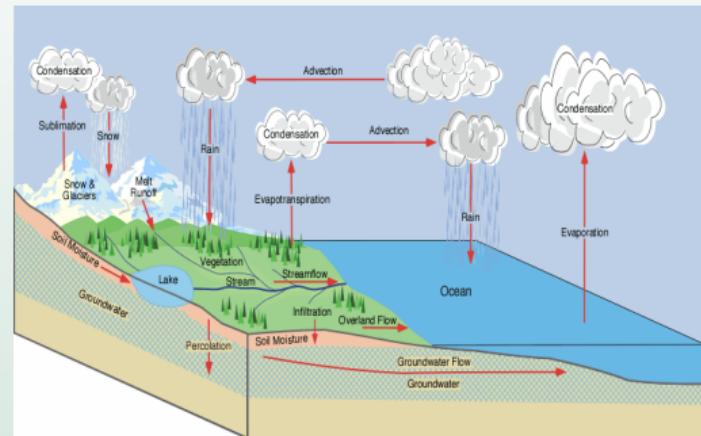


Pidwirny, M. (2006). "The Natural Spheres: The Hydrologic Cycle".

Fundamentals of Physical Geography, 2nd Edition

Hydrology

- The scientific study of the movement, distribution, and quality of water on Earth and other planets.
- Requires complicated physical models that are composed of many smaller components.



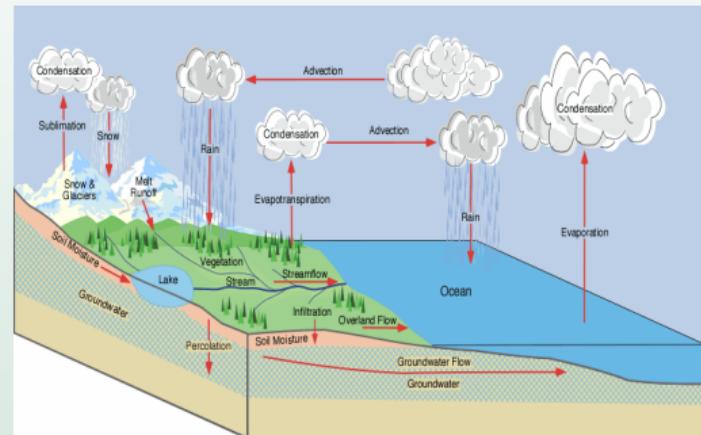
Pidwirny, M. (2006). "The Natural Spheres: The Hydrologic Cycle".

Fundamentals of Physical Geography, 2nd Edition

Hydrology

- The scientific study of the movement, distribution, and quality of water on Earth and other planets.
- Requires complicated physical models that are composed of many smaller components.

The science isn't always settled:



Pidwirny, M. (2006). "The Natural Spheres: The Hydrologic Cycle".

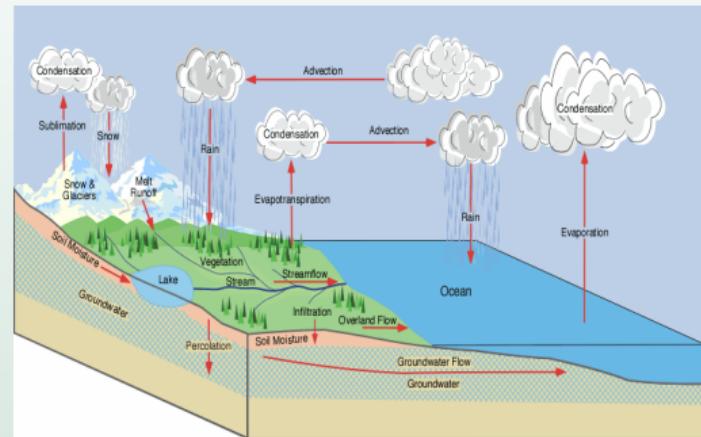
Fundamentals of Physical Geography, 2nd Edition

Hydrology

- The scientific study of the movement, distribution, and quality of water on Earth and other planets.
- Requires complicated physical models that are composed of many smaller components.

The science isn't always settled:

- How to model a given process?



Pidwirny, M. (2006). "The Natural Spheres: The Hydrologic Cycle".

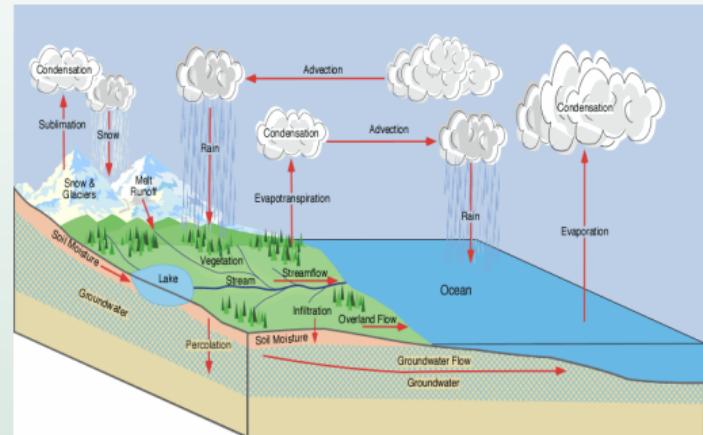
Fundamentals of Physical Geography, 2nd Edition

Hydrology

- The scientific study of the movement, distribution, and quality of water on Earth and other planets.
- Requires complicated physical models that are composed of many smaller components.

The science isn't always settled:

- How to model a given process?
- How do processes interact?



Pidwirny, M. (2006). "The Natural Spheres: The Hydrologic Cycle".

Fundamentals of Physical Geography, 2nd Edition

What is CHM?

The Canadian hydrological model (CHM) is a novel unstructured mesh-based approach to hydrological modelling. It is designed with a plug-in architecture of modular process representations to assist with the development and testing of hydrological models. The combination of unstructured meshes with plug-in processes allows CHM to bridge spatial scales, temporal scales, and spatial extents.

<https://github.com/Chrismarsh/CHM>

CHM

CHM

- is concerned mostly with modelling surface and ground water hydrology.

CHM

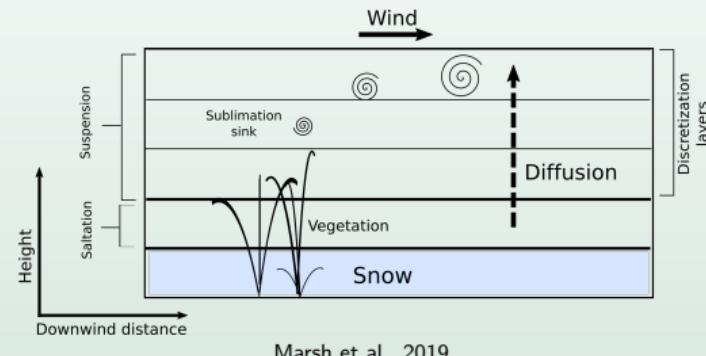
- is concerned mostly with modelling surface and ground water hydrology.
- is intended to handle spatially heterogeneous models required for snow redistribution phenomena (for example).

CHM

- is concerned mostly with modelling surface and ground water hydrology.
- is intended to handle spatially heterogeneous models required for snow redistribution phenomena (for example).
- is designed with the ability to plug in various *types* of other models as well.

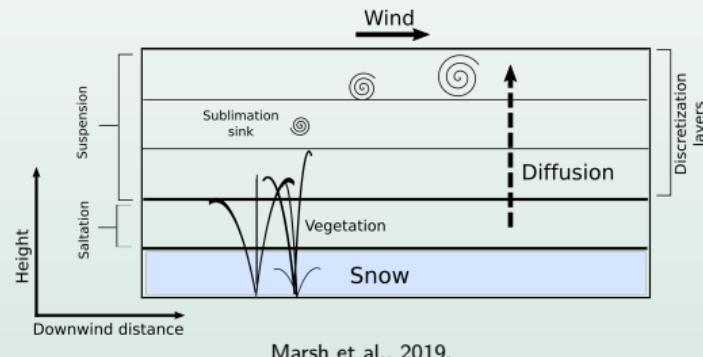
Example: Prairie blowing snow model (PBSM3D)

- *Partial differential equation* model for blowing snow.

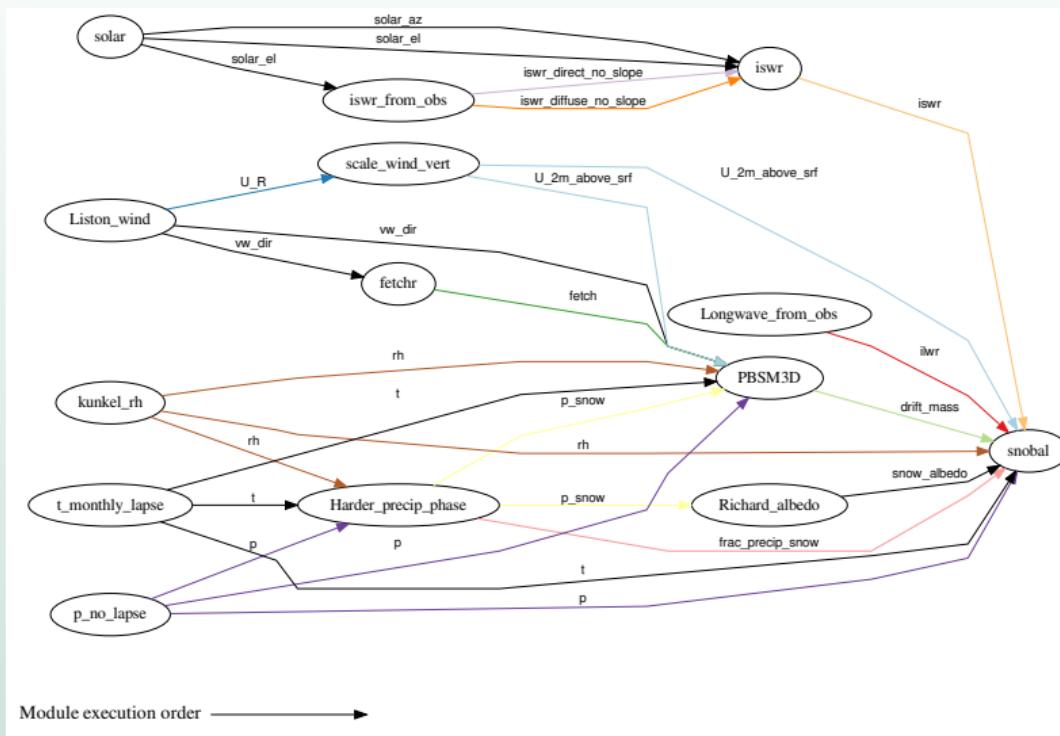


Example: Prairie blowing snow model (PBSM3D)

- *Partial differential equation* model for blowing snow.
- Saltation and suspension processes.
 - Multi-layer suspension.



A basic simulation with PBSM3D blowing snow



Space-dependence

- Some processes are influenced by (and influence) model behaviour in nearby regions of a domain.

Space-dependence

- Some processes are influenced by (and influence) model behaviour in nearby regions of a domain.
- Partial differential equations provide a means for describing such interaction with calculus. i.e.,

Space-dependence

- Some processes are influenced by (and influence) model behaviour in nearby regions of a domain.
- Partial differential equations provide a means for describing such interaction with calculus. i.e.,

$$\begin{aligned}-\nabla^2 u &= f(x, y) \\ -u_{xx}(x, y) - u_{yy}(x, y) &= f(x, y)\end{aligned}$$

Space-dependence

- Some processes are influenced by (and influence) model behaviour in nearby regions of a domain.
- Partial differential equations provide a means for describing such interaction with calculus. i.e.,

$$\begin{aligned}-\nabla^2 u &= f(x, y) \\ -u_{xx}(x, y) - u_{yy}(x, y) &= f(x, y)\end{aligned}$$

"Second derivative of u in the x direction plus the second derivative of u in the y direction equals a known function f "

Space-dependence

- Some processes are influenced by (and influence) model behaviour in nearby regions of a domain.
- Partial differential equations provide a means for describing such interaction with calculus. i.e.,

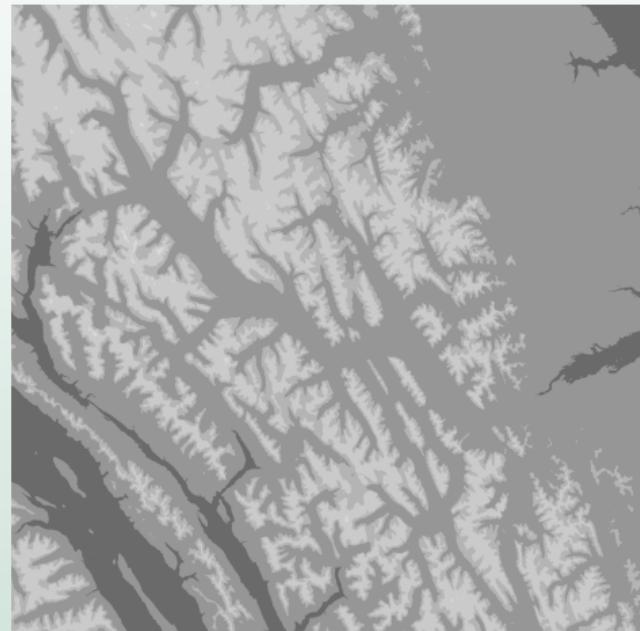
$$\begin{aligned}-\nabla^2 \textcolor{red}{u} &= f(x, y) \\ -u_{xx}(x, y) - u_{yy}(x, y) &= f(x, y)\end{aligned}$$

"Second derivative of u in the x direction plus the second derivative of u in the y direction equals a known function f "

→ What is $u(x, y)$ for all x and y ?

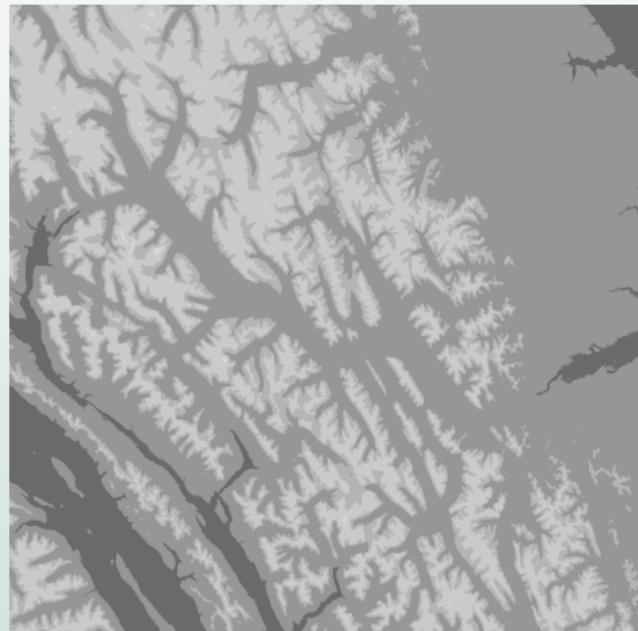
Spatial domains

- Geospatial data: latitude, longitude, elevation (*+ other data*)



Spatial domains

- Geospatial data: latitude, longitude, elevation (*+ other data*)
- Raster format (*data at each “pixel”*)



Mesh generation

Goal:

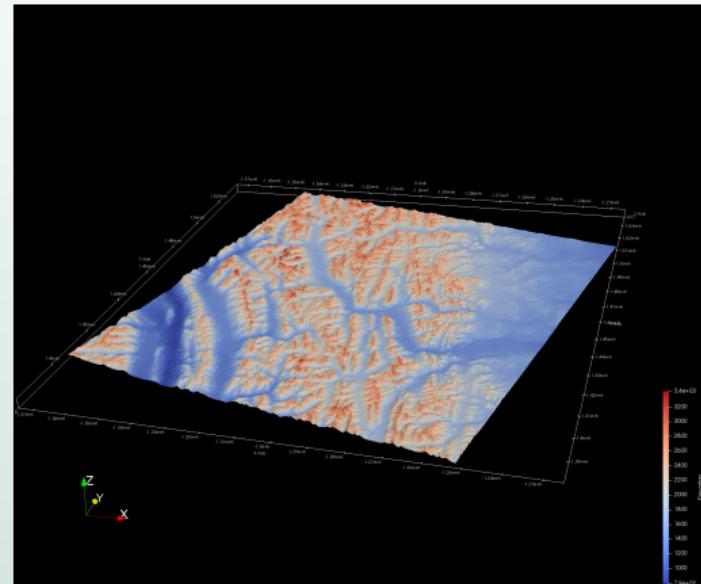
Produce a representation of the spatial domain

Mesh generation

Goal:

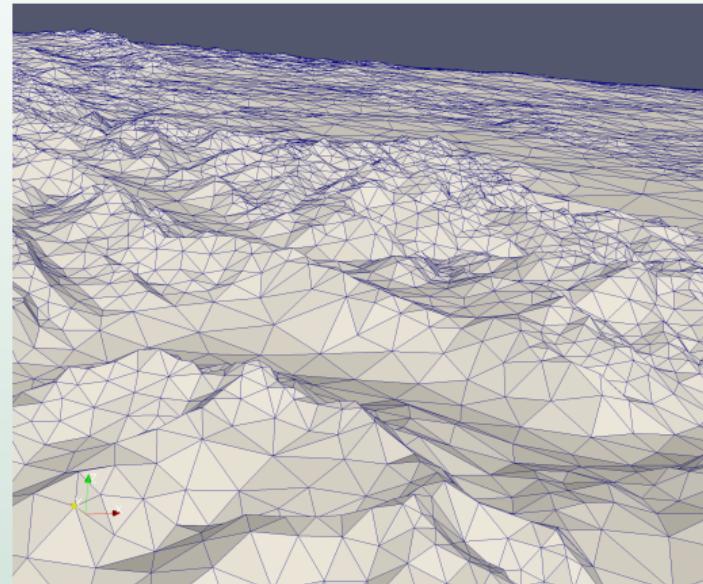
Produce a representation of the spatial domain

(This domain is $\sim 150 \text{ km} \times 150 \text{ km}$)



Triangular mesh

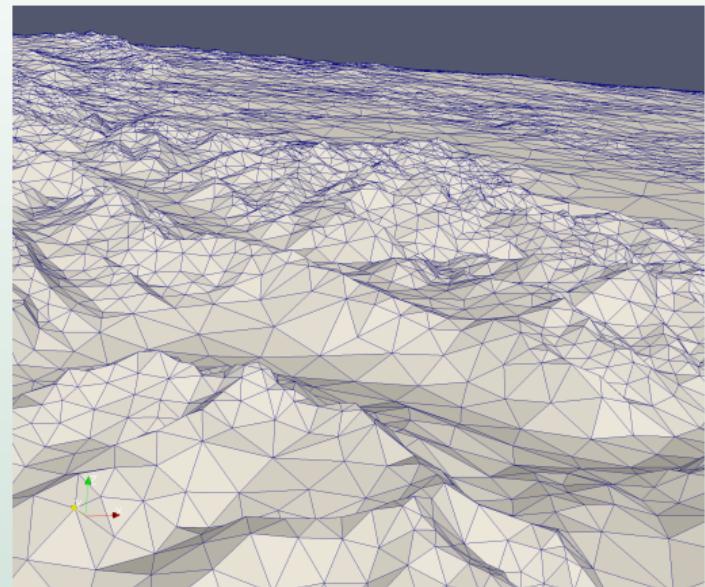
Spatial representation is made up of a bunch of planar triangles



Triangular mesh

Spatial representation is made up of a bunch of planar triangles

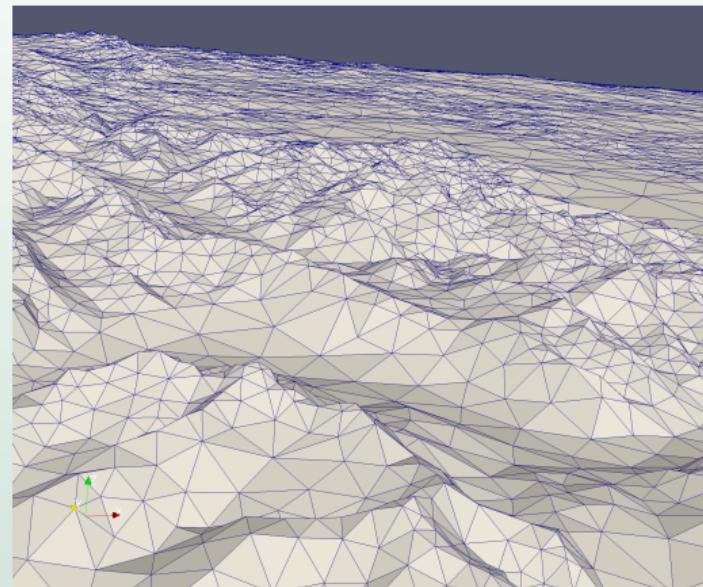
- variable size



Triangular mesh

Spatial representation is made up of a bunch of planar triangles

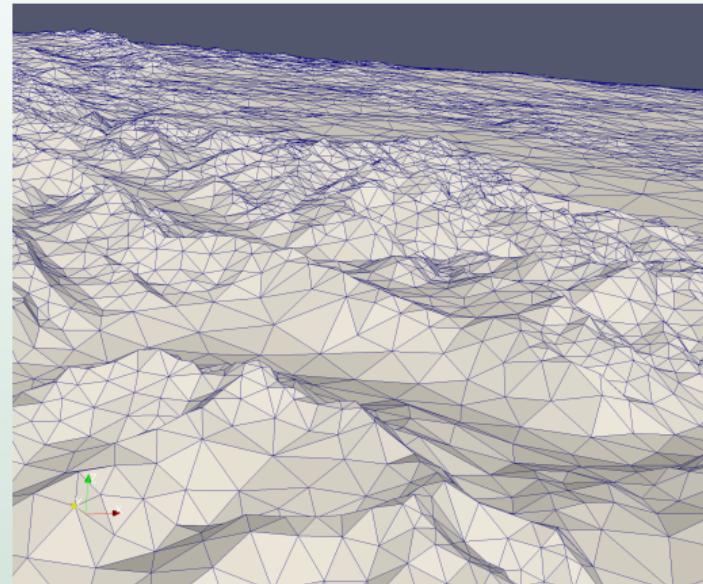
- variable size
- variable orientation



Triangular mesh

Spatial representation is made up of a bunch of planar triangles

- variable size
- variable orientation
- interpolating the input data

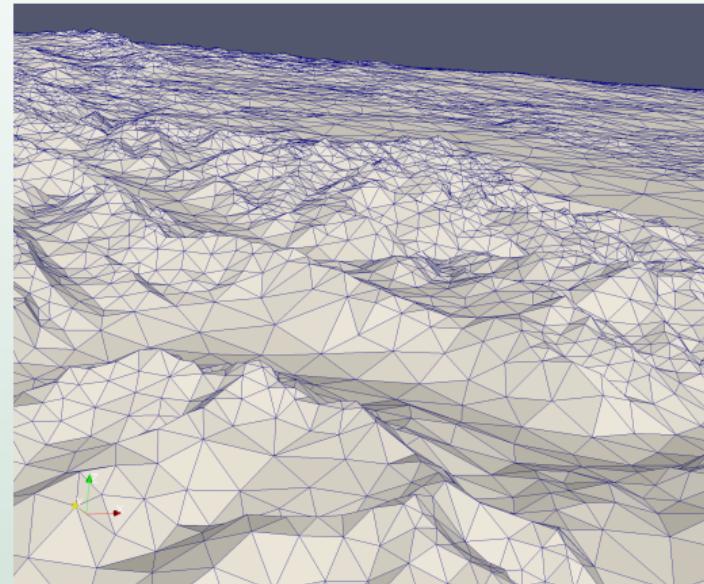


Triangular mesh

Spatial representation is made up of a bunch of planar triangles

- variable size
- variable orientation
- interpolating the input data

How do we get here?



Mesher

<https://github.com/Chrismarsh/mesher>

Mesher

Iteratively generate a triangular mesh according to various criteria, like elevation accuracy.

<https://github.com/Chrismarsh/mesher>

Mesher

Iteratively generate a triangular mesh according to various criteria, like elevation accuracy.

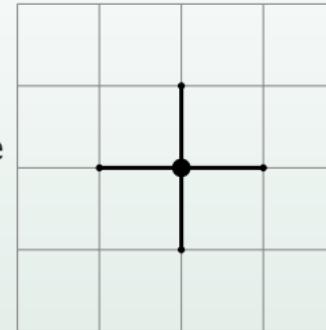
<https://github.com/Chrismarsh/mesher>

(Computational geometry. . . if you're interested)

Discretizing a PDE

Finite difference method

- approximate spatial derivatives with finite differences (only really useful for highly structured meshes)



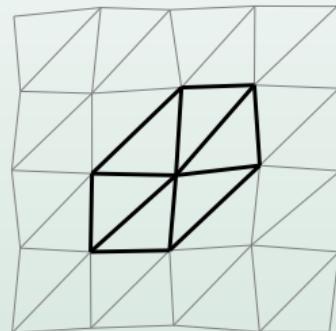
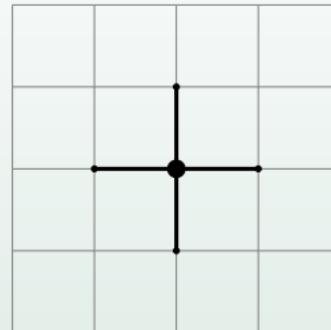
Discretizing a PDE

Finite difference method

- approximate spatial derivatives with finite differences (only really useful for highly structured meshes)

Finite element method

- using “nodal” points (vertices of mesh) to approximate function values over the “elements” (triangles)



Discretizing a PDE

Finite difference method

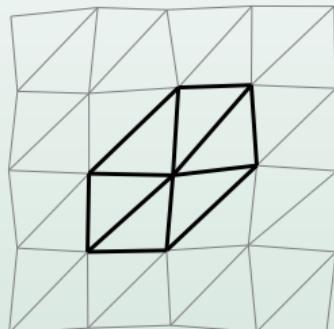
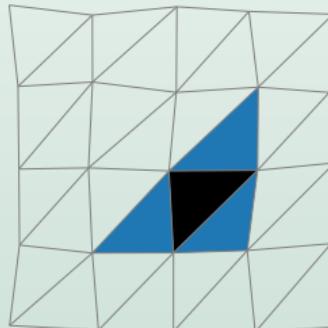
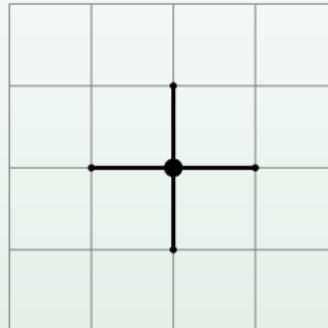
- approximate spatial derivatives with finite differences (only really useful for highly structured meshes)

Finite element method

- using “nodal” points (vertices of mesh) to approximate function values over the “elements” (triangles)

Finite volume method

- what flows out of one element (triangle), flows into its neighbour



Calculus to linear algebra

Calculus to linear algebra

Discretizing a linear PDE

Calculus to linear algebra

Discretizing a linear PDE

results in linear algebraic equations.

Calculus to linear algebra

Discretizing a linear PDE

results in linear algebraic equations.

There are many ways to do this.

Calculus to linear algebra

Discretizing a linear PDE

results in linear algebraic equations.

There are many ways to do this.

The mesh-based approaches listed above result in sparsely coupled equations.

Calculus to linear algebra

Discretizing a linear PDE

results in linear algebraic equations.

There are many ways to do this.

The mesh-based approaches listed above result in sparsely coupled equations.

(Numerical PDEs... if you're interested)

Sparse linear algebra

Solving

$$Au = b$$

Sparse linear algebra

Solving

$$\textcolor{red}{A}u = b$$

- $\textcolor{red}{A}$ - sparse matrix, discrete version of spatial influence

Sparse linear algebra

Solving

$$Au = \mathbf{b}$$

- A - sparse matrix, discrete version of spatial influence
- \mathbf{b} - known vector, discrete version of $f(x, y)$

Sparse linear algebra

Solving

$$A\mathbf{u} = \mathbf{b}$$

- A - sparse matrix, discrete version of spatial influence
- \mathbf{b} - known vector, discrete version of $f(x, y)$
- \mathbf{u} - unknown vector, discrete version of $u(x, y)$

Sparse linear algebra

Solving

$$Au = b$$

There are many ways to do this too.

- A - sparse matrix, discrete version of spatial influence
- b - known vector, discrete version of $f(x, y)$
- u - unknown vector, discrete version of $u(x, y)$

Direct methods

Direct “inversion” of linear systems.

- Obtain numerically precise solution
 $|Au - b| \sim \mathcal{O}(\epsilon_{\text{mach}})$

ϵ_{mach} – machine epsilon

Direct methods

Direct “inversion” of linear systems.

- Obtain numerically precise solution
 $|Au - b| \sim \mathcal{O}(\epsilon_{\text{mach}})$

Methods:

- Cholesky

[Benoit1924]

$$PAP^T = LL^T$$

ϵ_{mach} – machine epsilon

Direct methods

Direct “inversion” of linear systems.

- Obtain numerically precise solution
 $|Au - b| \sim \mathcal{O}(\epsilon_{\text{mach}})$

ϵ_{mach} – machine epsilon

Methods:

- Cholesky

[Benoit1924]

$$PAP^T = LL^T$$

- LU

[Banachiewicz1938]

$$PAQ^T = LU$$

Direct methods

Direct “inversion” of linear systems.

- Obtain numerically precise solution
 $|Au - b| \sim \mathcal{O}(\epsilon_{\text{mach}})$

ϵ_{mach} – machine epsilon

P, Q – permutation matrices

Methods:

- Cholesky

[Benoit1924]

$$PAP^T = LL^T$$

- LU

[Banachiewicz1938]

$$PAQ^T = LU$$

Pros and cons of direct methods

Pros:

Cons:

Pros and cons of direct methods

Pros:

- As accurate as you can get numerically

Cons:

Pros and cons of direct methods

Pros:

- As accurate as you can get numerically

Cons:

- Algorithms scale poorly ($\mathcal{O}(m^3)$ in m unknowns, for a general sparse matrix)

Pros and cons of direct methods

Pros:

- As accurate as you can get numerically
- Re-use factorization to solve with different b

Cons:

- Algorithms scale poorly ($\mathcal{O}(m^3)$ in m unknowns, for a general sparse matrix)

Pros and cons of direct methods

Pros:

- As accurate as you can get numerically
- Re-use factorization to solve with different b

Cons:

- Algorithms scale poorly ($\mathcal{O}(m^3)$ in m unknowns, for a general sparse matrix)

Pro/Con:

- Time to solution dependent *entirely* on matrix structure
 - *might* be substantially better than worst-case

Permutations

Free to rearrange our equations and unknowns.

Permutations

Free to rearrange our equations and unknowns.

Within reason:

- Numerical stability considerations

Permutations

Free to rearrange our equations and unknowns.

Within reason:

- Numerical stability considerations

Can be used to alter matrix structure...

Permutations

Free to rearrange our equations and unknowns.

$$Au = b$$

Within reason:

- Numerical stability considerations

Can be used to alter matrix structure...

view as matrix multiplication:

Permutations

Free to rearrange our equations and unknowns.

Within reason:

- Numerical stability considerations

Can be used to alter matrix structure...

view as matrix multiplication:

$$Au = b$$

$$\textcolor{red}{P}Au = \textcolor{red}{P}b$$

Permutations

Free to rearrange our equations and unknowns.

Within reason:

- Numerical stability considerations

Can be used to alter matrix structure...

view as matrix multiplication:

$$Au = b$$

$$PAu = Pb$$

$$PA(Q^T Q)u = Pb$$

Permutations

Free to rearrange our equations and unknowns.

Within reason:

- Numerical stability considerations

Can be used to alter matrix structure...

view as matrix multiplication:

$$Au = b$$

$$PAu = Pb$$

$$PA(Q^T Q)u = Pb$$

$$(PAQ^T)(Qu) = Pb$$

Permutations

Free to rearrange our equations and unknowns.

Within reason:

- Numerical stability considerations

Can be used to alter matrix structure...

view as matrix multiplication:

$$Au = b$$

$$PAu = Pb$$

$$PA(Q^T Q)u = Pb$$

$$(PAQ^T)(Qu) = Pb$$

$$\tilde{A}\tilde{u} = \tilde{b}$$

Permutations

Free to rearrange our equations and unknowns.

Within reason:

- Numerical stability considerations

Can be used to alter matrix structure...

view as matrix multiplication:

$$Au = b$$

$$PAu = Pb$$

$$PA(Q^T Q)u = Pb$$

$$(PAQ^T)(Qu) = Pb$$

$$\tilde{A}\tilde{u} = \tilde{b}$$

Permutations

Free to rearrange our equations and unknowns.

Within reason:

- Numerical stability considerations

Can be used to alter matrix structure...

view as matrix multiplication:

$$Au = b$$

$$PAu = Pb$$

$$PA(Q^T Q)u = Pb$$

$$(PAQ^T)(Qu) = \textcolor{red}{Pb}$$

$$\tilde{A}\tilde{u} = \tilde{b}$$

Permutations

Free to rearrange our equations and unknowns.

Within reason:

- Numerical stability considerations

Can be used to alter matrix structure...

view as matrix multiplication:

$$Au = b$$

$$PAu = Pb$$

$$PA(Q^T Q)u = Pb$$

$$(PAQ^T)(Qu) = Pb$$

$$\tilde{A}\tilde{u} = \tilde{b}$$

Want \tilde{A} to have a *better* structure.

Symmetric permutations

Changing equation *and* unknown ordering simultaneously.

Symmetric permutations

Changing equation *and* unknown ordering simultaneously.

- Permutation matrix P

Symmetric permutations

Changing equation *and* unknown ordering simultaneously.

- Permutation matrix P
- $A \rightarrow PAP^T$

Symmetric permutations

Changing equation *and* unknown ordering simultaneously.

- Permutation matrix P
- $A \rightarrow PAP^T$

Nice properties:

Symmetric permutations

Changing equation *and* unknown ordering simultaneously.

- Permutation matrix P
- $A \rightarrow PAP^T$

Nice properties:

- If A is symmetric positive definite (SPD), so is PAP^T .

Symmetric permutations

Changing equation *and* unknown ordering simultaneously.

- Permutation matrix P
- $A \rightarrow PAP^T$

Nice properties:

- If A is symmetric positive definite (SPD), so is PAP^T .
- If A is diagonally dominant, so is PAP^T .

Symmetric permutations

Changing equation *and* unknown ordering simultaneously.

- Permutation matrix P
- $A \rightarrow PAP^T$

Nice properties:

- If A is symmetric positive definite (**SPD**), so is PAP^T .
 - If A is diagonally dominant, so is PAP^T .
- No loss of stability for certain PDE operators.

Bandwidth-reducing order

Goal: Minimize bandwidth of A .

Bandwidth-reducing order

Goal: Minimize bandwidth of A .

Why?

Bandwidth-reducing order

Goal: Minimize bandwidth of A .

Why?

- Bandwidth of factors is less than or equal to bandwidth of original matrix

Bandwidth-reducing order

Goal: Minimize bandwidth of A .

Why?

- Bandwidth of factors is less than or equal to bandwidth of original matrix
- Less chance for “fill-in” to occur in factors.

Bandwidth-reducing order

Goal: Minimize bandwidth of A .

Why?

- Bandwidth of factors is less than or equal to bandwidth of original matrix
- Less chance for “fill-in” to occur in factors.
- “Banded” solvers can exploit this structure for efficiency.

Bandwidth-reducing order

Goal: Minimize bandwidth of A .

Why?

- Bandwidth of factors is less than or equal to bandwidth of original matrix
- Less chance for “fill-in” to occur in factors.
- “Banded” solvers can exploit this structure for efficiency.
 - LAPACK routines like `dgbtr{fs}`

Bandwidth-reducing order

Goal: Minimize bandwidth of A .

Why?

- Bandwidth of factors is less than or equal to bandwidth of original matrix
- Less chance for “fill-in” to occur in factors.
- “Banded” solvers can exploit this structure for efficiency.
 - LAPACK routines like `dgbtr{fs}` → **Double precision**, general banded, triangular, factor/solve.

Bandwidth-reducing order

Goal: Minimize bandwidth of A .

Why?

- Bandwidth of factors is less than or equal to bandwidth of original matrix
- Less chance for “fill-in” to occur in factors.
- “Banded” solvers can exploit this structure for efficiency.
 - LAPACK routines like `dgbtr{fs}` → Double precision, general banded, triangular, factor/solve.

Bandwidth-reducing order

Goal: Minimize bandwidth of A .

Why?

- Bandwidth of factors is less than or equal to bandwidth of original matrix
- Less chance for “fill-in” to occur in factors.
- “Banded” solvers can exploit this structure for efficiency.
 - LAPACK routines like `dgbtr{fs}` → Double precision, general banded, `triangular`, factor/solve.

Bandwidth-reducing order

Goal: Minimize bandwidth of A .

Why?

- Bandwidth of factors is less than or equal to bandwidth of original matrix
- Less chance for “fill-in” to occur in factors.
- “Banded” solvers can exploit this structure for efficiency.
 - LAPACK routines like `dgbtr{fs}` → Double precision, general banded, triangular, `factor/solve`.

Bandwidth-reducing order

Goal: **Minimize** bandwidth of A .

How? (*NP-Hard*) problem

Why?

- Bandwidth of factors is less than or equal to bandwidth of original matrix
- Less chance for “fill-in” to occur in factors.
- “Banded” solvers can exploit this structure for efficiency.
 - LAPACK routines like `dgbtr{fs}` → Double precision, general banded, triangular, factor/solve.

Bandwidth-reducing order

Goal: Minimize bandwidth of A .

How? (*NP-Hard*) problem

Why?

- Bandwidth of factors is less than or equal to bandwidth of original matrix
- Less chance for “fill-in” to occur in factors.
- “Banded” solvers can exploit this structure for efficiency.
 - LAPACK routines like `dgbtr{fs}` → Double precision, general banded, triangular, factor/solve.

Heuristic algorithms

- Cuthill–McKee (CM), [Cuthill1969]

Bandwidth-reducing order

Goal: Minimize bandwidth of A .

Why?

- Bandwidth of factors is less than or equal to bandwidth of original matrix
- Less chance for “fill-in” to occur in factors.
- “Banded” solvers can exploit this structure for efficiency.
 - LAPACK routines like `dgbtr{fs}` → Double precision, general banded, triangular, factor/solve.

How? (*NP-Hard*) problem

Heuristic algorithms

- Cuthill–McKee (CM), [Cuthill1969]
- Reverse Cuthill–McKee (RCM), [George1971]

Bandwidth-reducing order

Goal: Minimize bandwidth of A .

Why?

- Bandwidth of factors is less than or equal to bandwidth of original matrix
- Less chance for “fill-in” to occur in factors.
- “Banded” solvers can exploit this structure for efficiency.
 - LAPACK routines like `dgbtr{fs}` → Double precision, general banded, triangular, factor/solve.

How? (*NP-Hard*) problem

Heuristic algorithms

- Cuthill–McKee (CM), [Cuthill1969]
- Reverse Cuthill–McKee (RCM), [George1971]

We use **RCM** for reduced bandwidth order.

Fill-reducing ordering

Goal: Minimize amount of fill-in in factors. (*Also NP-Hard*)

Fill-reducing ordering

Goal: Minimize amount of fill-in in factors. (*Also NP-Hard*)

Why?

Fill-reducing ordering

Goal: Minimize amount of fill-in in factors. (*Also NP-Hard*)

Why?

- Sparse forward/backward substitution only operates on entries that are there.
 - $\mathcal{O}(nnz)$, nnz - number of non-zeros

Fill-reducing ordering

Goal: Minimize amount of fill-in in factors. (*Also NP-Hard*)

Why?

- Sparse forward/backward substitution only operates on entries that are there.
 - $\mathcal{O}(nnz)$, nnz - number of non-zeros
- Fewer operations to perform in forward/backward substitution.

Fill-reducing ordering

Goal: Minimize amount of fill-in in factors. (*Also NP-Hard*)

Why?

- Sparse forward/backward substitution only operates on entries that are there.
 - $\mathcal{O}(nnz)$, nnz - number of non-zeros
- Fewer operations to perform in forward/backward substitution.

How? (*Heuristics, again*)

Fill-reducing ordering

Goal: Minimize amount of fill-in in factors. (*Also NP-Hard*)

Why?

- Sparse forward/backward substitution only operates on entries that are there.
 - $\mathcal{O}(nnz)$, nnz - number of non-zeros
- Fewer operations to perform in forward/backward substitution.

How? (*Heuristics, again*)

- Algorithms based on Multilevel Nested Dissection (ND) [George1973]

Specific problem (from PBSM3D)

For stability of time-integration, requires solution to a steady-state PDE at each time step:

$$(\mathbb{I} - D\nabla^2) u = f(x, y)$$

Specific problem (from PBSM3D)

For stability of time-integration, requires solution to a steady-state PDE at each time step:

$$(\mathbb{I} - D\nabla^2) u = f(x, y)$$

- D – smoothing coefficient
- $u(x, y)$ – rate of snow mass transport at x, y coordinate
- $f(x, y)$ – known quantity dependent on other physical parameters

Specific problem (from PBSM3D)

For stability of time-integration, requires solution to a steady-state PDE at each time step:

$$(\mathbb{I} - D\nabla^2) u = f(x, y)$$

(Requires much more, but we're just going to stick to this)

- D – smoothing coefficient
- $u(x, y)$ – rate of snow mass transport at x, y coordinate
- $f(x, y)$ – known quantity dependent on other physical parameters

Discretization of specific problem

$$Au = b$$

Discretization of specific problem

$$Au = b$$

A represents a finite volume discretization of $(\mathbb{I} - D\nabla^2)$.
(A is symmetric positive definite.)

Discretization of specific problem

$$Au = b$$

A represents a finite volume discretization of $(\mathbb{I} - D\nabla^2)$.
(A is symmetric positive definite.)

Components of u represent the rate of snow mass transport, constant on each element.

Discretization of specific problem

$$Au = b$$

A represents a finite volume discretization of $(\mathbb{I} - D\nabla^2)$.
(A is symmetric positive definite.)

Components of u represent the rate of snow mass transport, constant on each element.

Components of b represent known quantities dependent on current snow mass (among other things), constant on each element.

Discretization of specific problem

$$Au = b$$

A represents a finite volume discretization of $(\mathbb{I} - D\nabla^2)$.
(A is symmetric positive definite.)

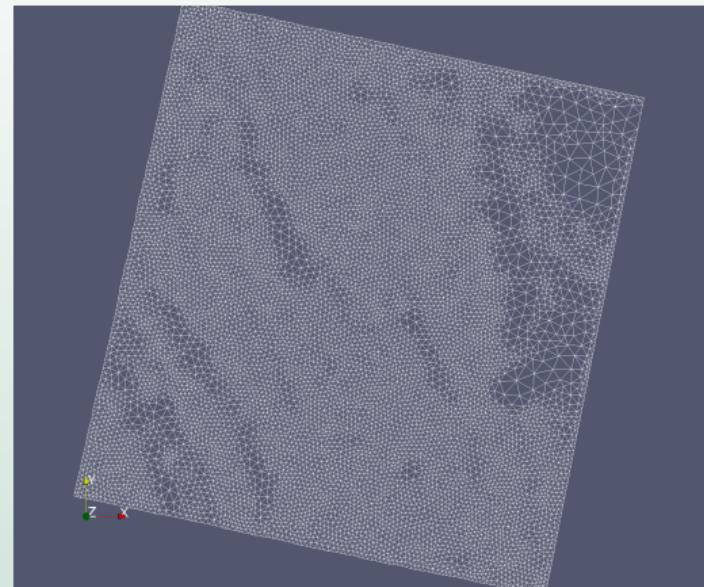
Components of u represent the rate of snow mass transport, constant on each element.

Components of b represent known quantities dependent on current snow mass (among other things), constant on each element.

Use known solution $u^* = 1^{N \times 1}$, $b^* = Au^*$ for testing,
i.e., solve $Au = b^*$.

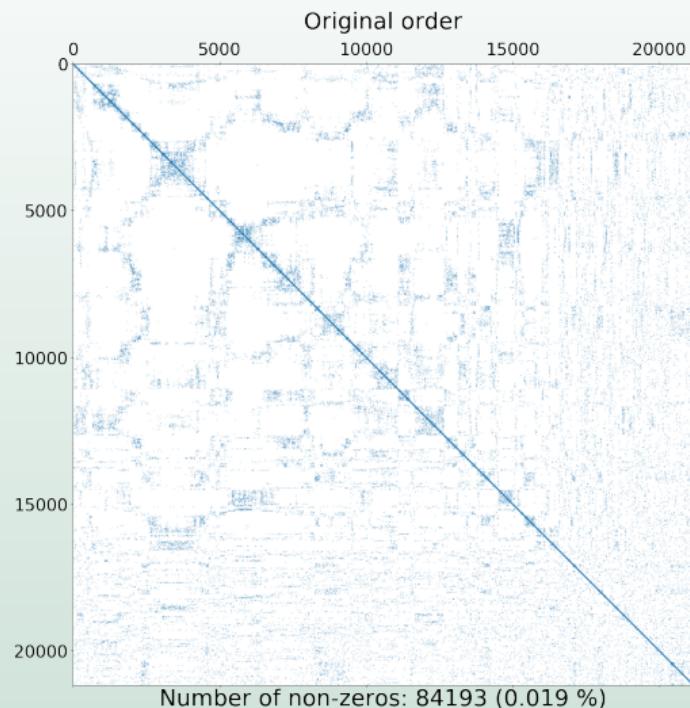
Mesh: coarse resolution

Low resolution (21168 elements)



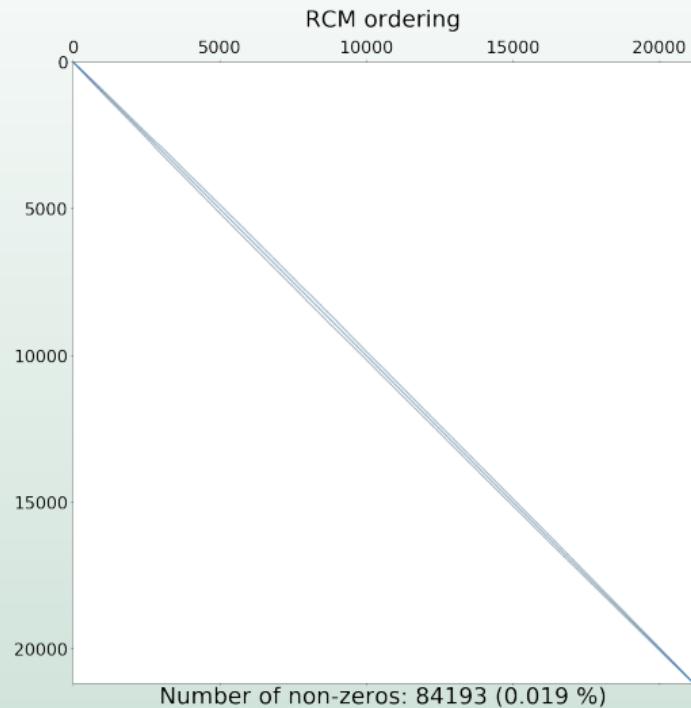
Matrix structure: original order

Bandwidth: 21062



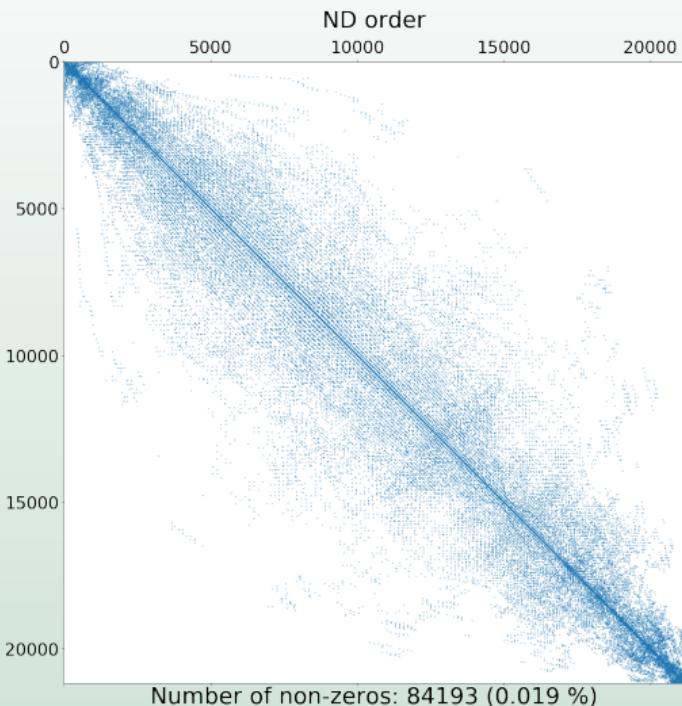
Matrix structure: RCM order

Bandwidth: 168



Matrix structure: ND order

Bandwidth: 12269



Performance: coarse mesh

Not shown here.

Performance: coarse mesh

Not shown here.

Jupyter notebook

- <https://github.com/uofs-simlab/Math313>
- Can run easily from colab.research.google.com if you have a Google account.

Performance: coarse mesh

Not shown here.

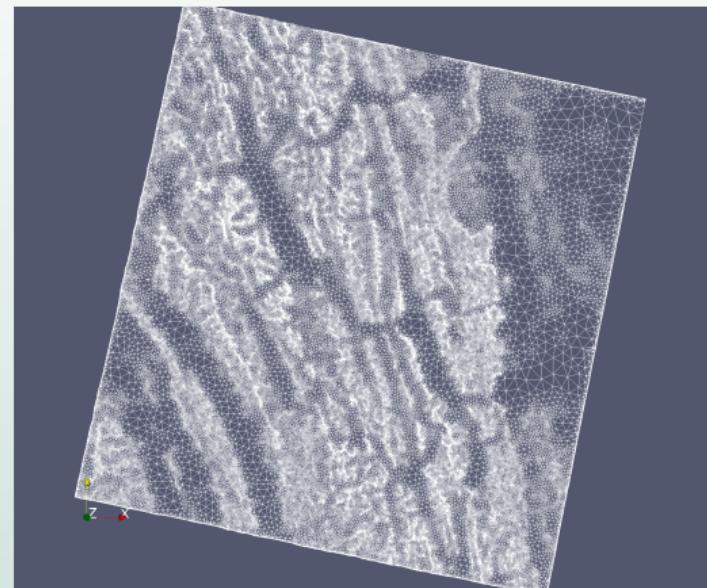
Jupyter notebook

- <https://github.com/uofs-simlab/Math313>
- Can run easily from colab.research.google.com if you have a Google account.

I went bigger...

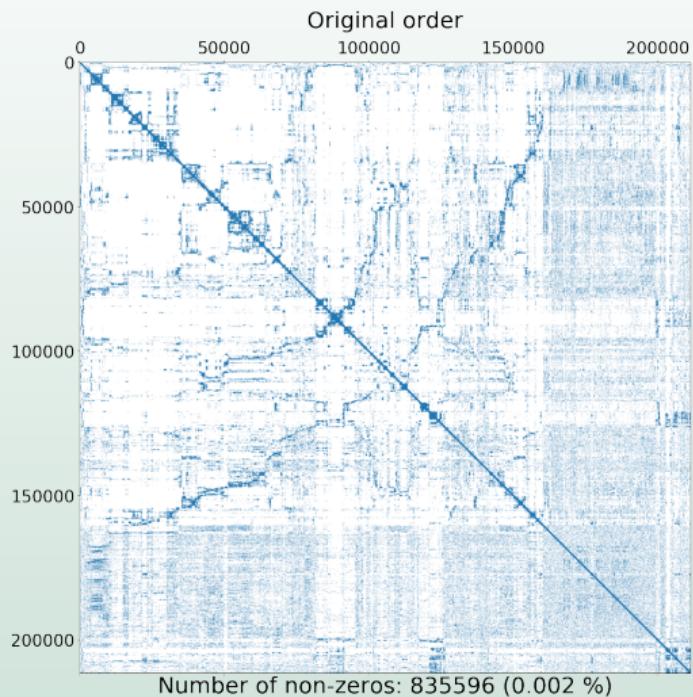
Mesh: fine resolution

High resolution (211415 elements)



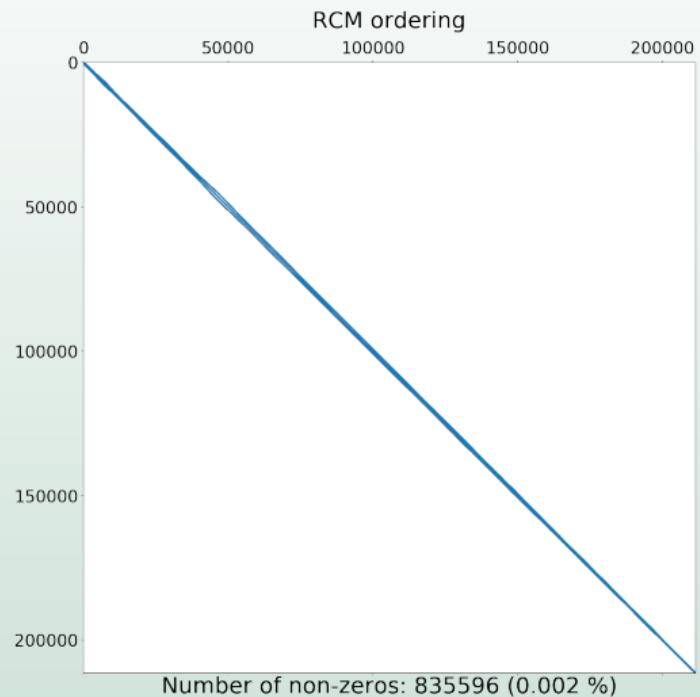
Matrix structure: original order

Bandwidth: 210855



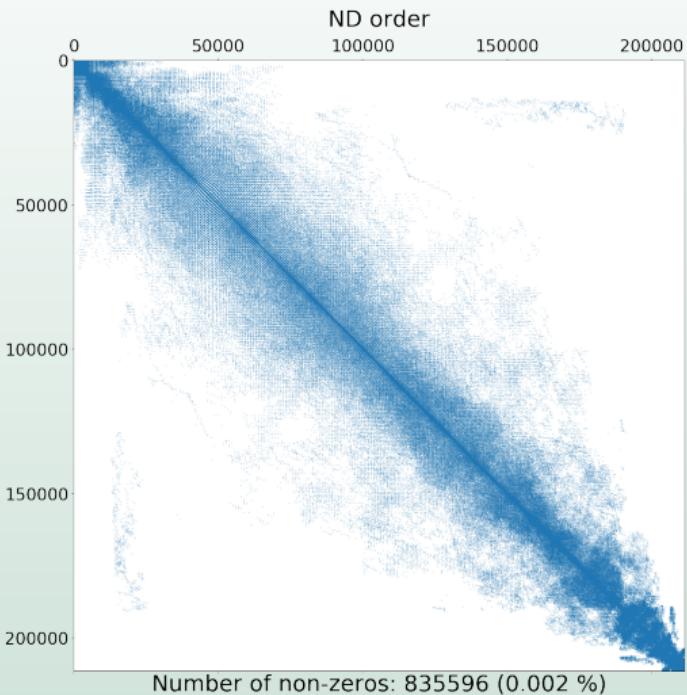
Matrix structure: RCM order

Bandwidth: 1320



Matrix: ND order

Bandwidth: 173136



Performance: LU – Original order

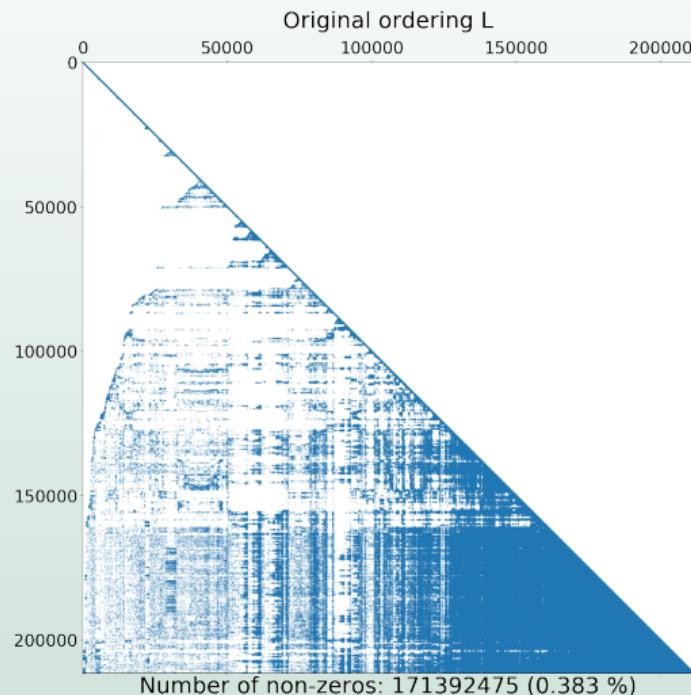
LU factorization

CPU times:

user: 8d 12h 41min 15s,
sys: 18h 1min 11s,
total: 9d 6h 42min 26s

Wall time: 6h 58min 57s

Solve



Performance: LU – Original order

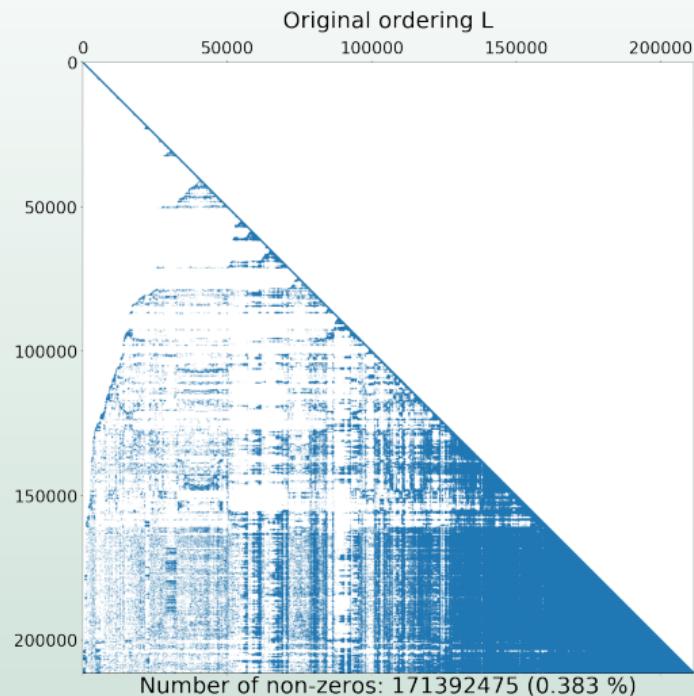
LU factorization

CPU times:

user: 8d 12h 41min 15s,
sys: 18h 1min 11s,
total: 9d 6h 42min 26s

Wall time: **6h 58min 57s**

Solve



Performance: LU – Original order

LU factorization

CPU times:

user: 8d 12h 41min 15s,
sys: 18h 1min 11s,
total: 9d 6h 42min 26s

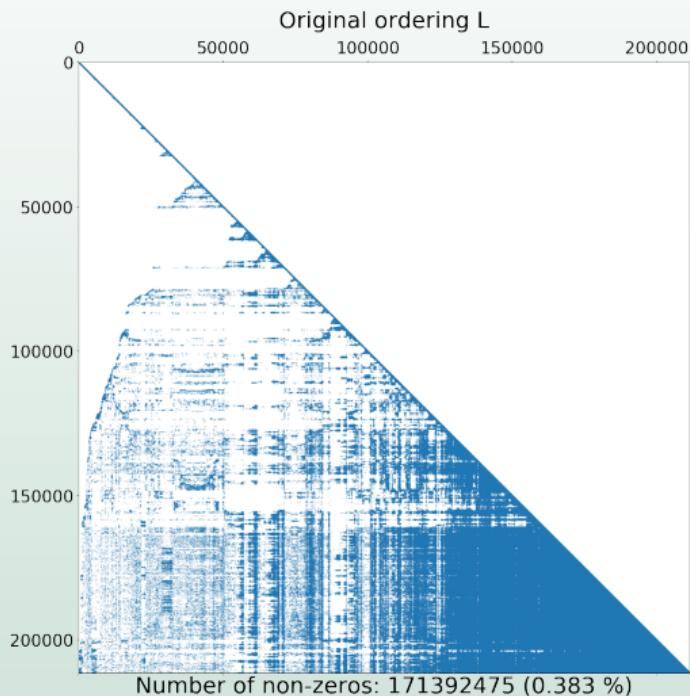
Wall time: 6h 58min 57s

Solve

Wall time: **8.74s**

Total

Wall time: 6h 59min 5s



Performance: LU – Original order

LU factorization

CPU times:

user: 8d 12h 41min 15s,
sys: 18h 1min 11s,
total: 9d 6h 42min 26s

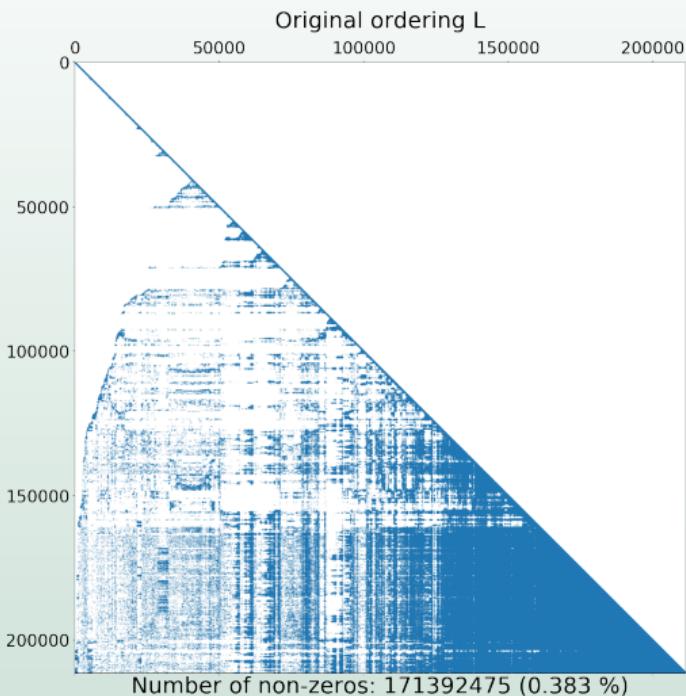
Wall time: 6h 58min 57s

Solve

Wall time: 8.74s

Total

Wall time: 6h 59min 5s



Performance: LU – RCM order

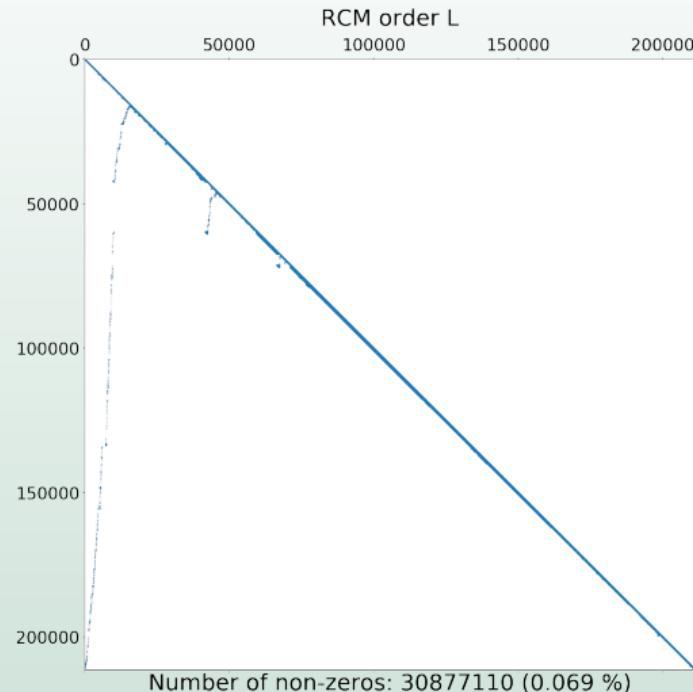
LU factorization

CPU times:

user: 56.2 s,
sys: 2.71 s,
total: 58.9 s

Wall time: 52.6 s

Solve



Performance: LU – RCM order

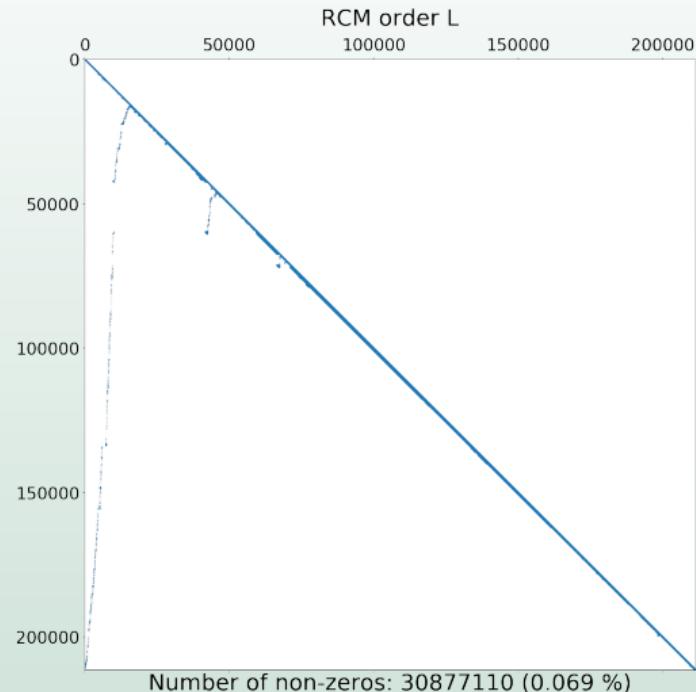
LU factorization

CPU times:

user: 56.2 s,
sys: 2.71 s,
total: 58.9 s

Wall time: 52.6 s

Solve



Performance: LU – RCM order

LU factorization

CPU times:

user: 56.2 s,
sys: 2.71 s,
total: 58.9 s

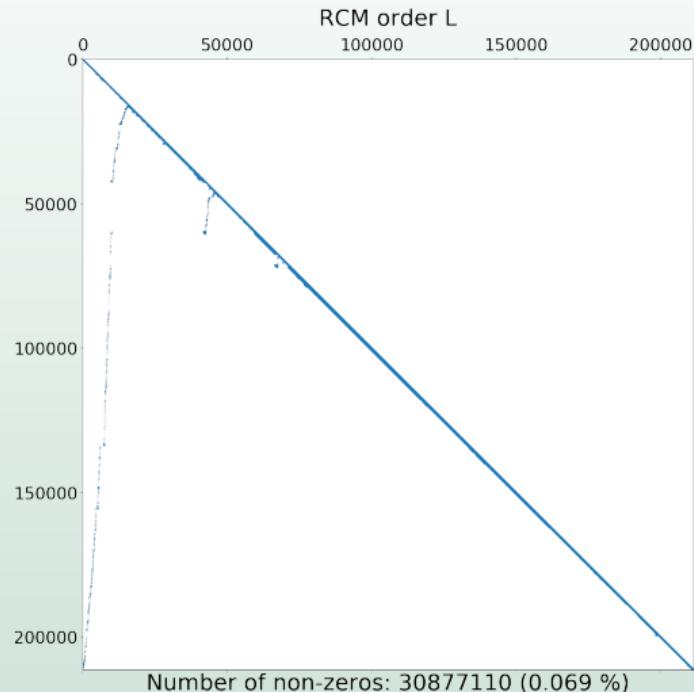
Wall time: 52.6 s

Solve

Wall time: 0.741 s

Total

Wall time: 53.3s



Performance: LU – RCM order

LU factorization

CPU times:

user: 56.2 s,
sys: 2.71 s,
total: 58.9 s

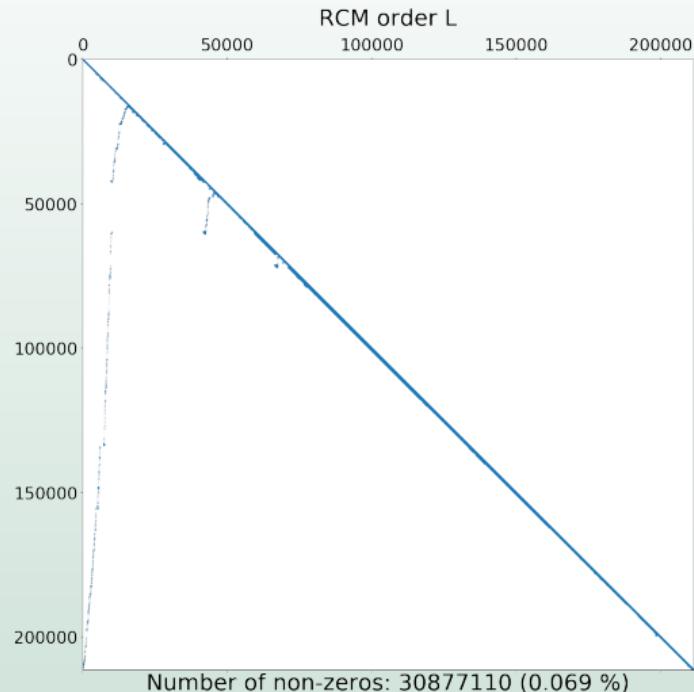
Wall time: 52.6 s

Solve

Wall time: 0.741 s

Total

Wall time: 53.3s



Performance: LU – ND order

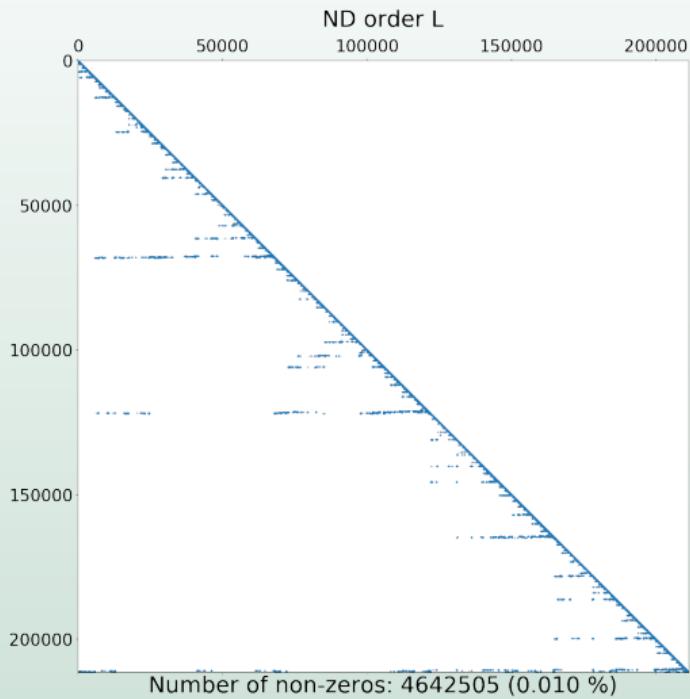
LU factorization

CPU times:

user: 6.67 s,
sys: 0.021 s,
total: 6.67 s

Wall time: 1.25 s

Solve



Performance: LU – ND order

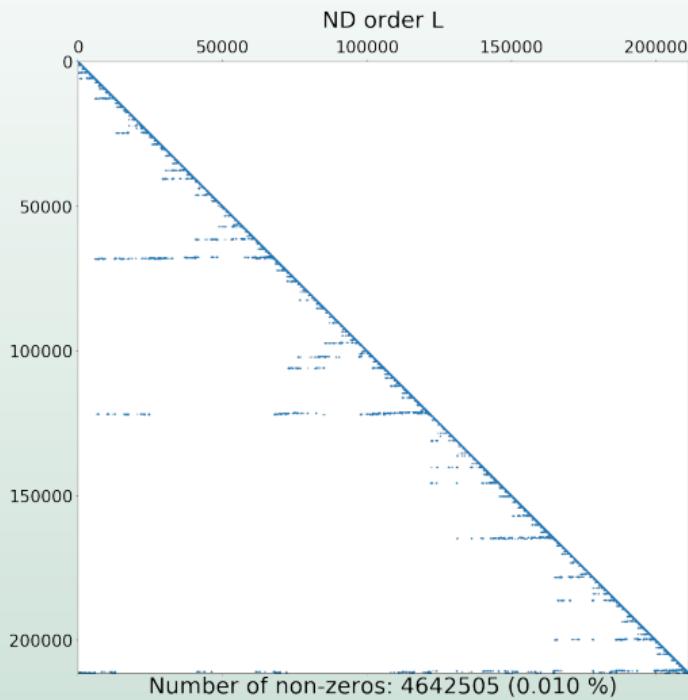
LU factorization

CPU times:

user: 6.67 s,
sys: 0.021 s,
total: 6.67 s

Wall time: 1.25 s

Solve



Performance: LU – ND order

LU factorization

CPU times:

user: 6.67 s,
sys: 0.021 s,
total: 6.67 s

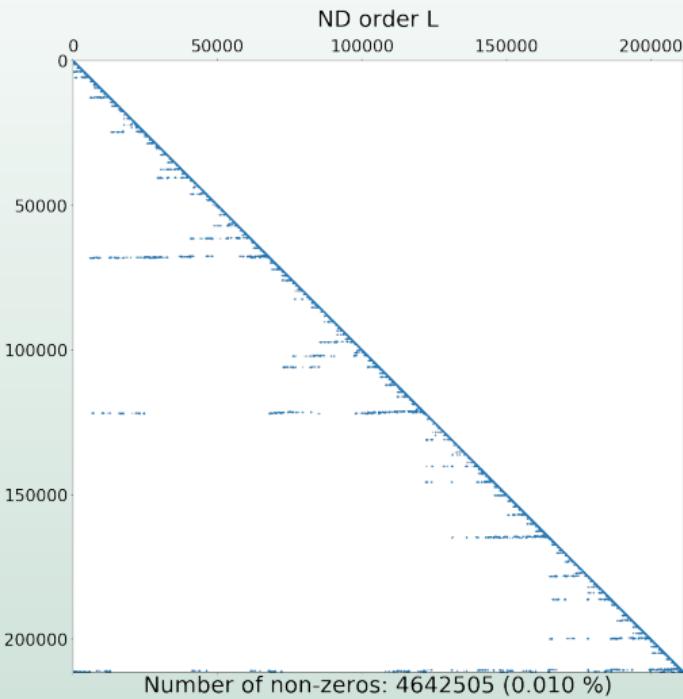
Wall time: 1.25 s

Solve

Wall time: 0.049 s

Total

Wall time: 1.30s



Performance: LU – ND order

LU factorization

CPU times:

user: 6.67 s,
sys: 0.021 s,
total: 6.67 s

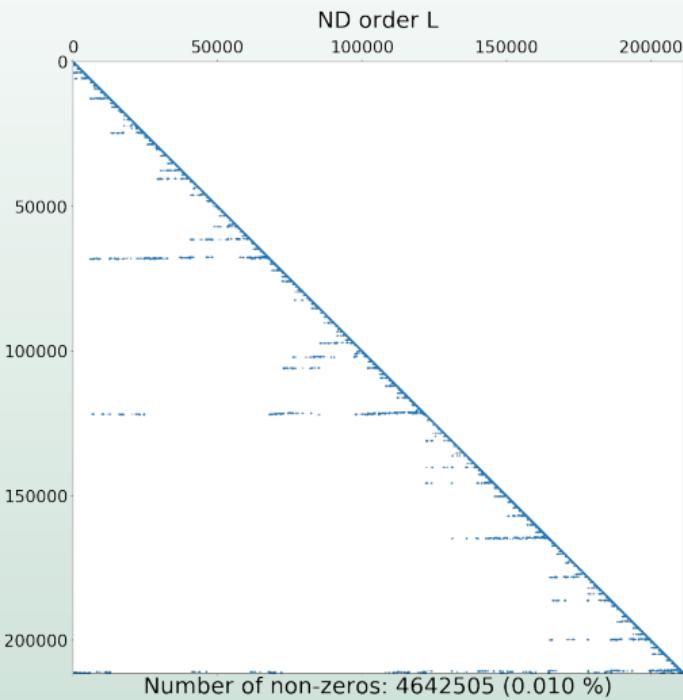
Wall time: 1.25 s

Solve

Wall time: 0.049 s

Total

Wall time: **1.30s**



Fine mesh summary

211415 elements → 835596 non-zero matrix entries

LU		
Order	Time (s)	Accuracy
original	25145	1.2e-13
RCM	53.3	1.6e-13
ND	1.3	1.5e-13

Concluding remarks

- A lot of pieces have to come together in the real world before we get to linear algebra.

Concluding remarks

- A lot of pieces have to come together in the real world before we get to linear algebra.
- There's still a lot to investigate once we get there.

Concluding remarks

- A lot of pieces have to come together in the real world before we get to linear algebra.
- There's still a lot to investigate once we get there.
- Substantial performance improvement observed in direct solve simply by reordering.

Concluding remarks

- A lot of pieces have to come together in the real world before we get to linear algebra.
- There's still a lot to investigate once we get there.
- Substantial performance improvement observed in direct solve simply by reordering.
- We haven't even talked about:

Concluding remarks

- A lot of pieces have to come together in the real world before we get to linear algebra.
- There's still a lot to investigate once we get there.
- Substantial performance improvement observed in direct solve simply by reordering.
- We haven't even talked about:
 - iterative algorithms (*next guest lecture!*)

Concluding remarks

- A lot of pieces have to come together in the real world before we get to linear algebra.
- There's still a lot to investigate once we get there.
- Substantial performance improvement observed in direct solve simply by reordering.
- We haven't even talked about:
 - iterative algorithms (*next guest lecture!*)
 - machine-dependent things (*High-performance computing... if you're interested*)