



## Numerical linear algebra in the Canadian hydrological model

### Part II: Iterative methods

Kevin R. Green

Department of Computer Science  
University of Saskatchewan

Math 313  
University of Saskatchewan  
October 13, 2020

# Outline

1 Review

2 Iterative solvers

3 Conclusions

## Brief review of last lecture

- Hydrological models consist of many sub-components.
  - Some involve spatial coupling via PDEs
- Discretizing PDEs on a mesh gives sparse linear algebra problems
- Time for sparse direct solves is sensitive to the sparsity structure

# Sparse linear algebra

## Solving

$$Au = b$$

# Sparse linear algebra

Solving

$$\textcolor{red}{A}u = b$$

- $\textcolor{red}{A}$  - sparse matrix, discrete version of spatial influence

# Sparse linear algebra

Solving

$$Au = \mathbf{b}$$

- $A$  - sparse matrix, discrete version of spatial influence
- $\mathbf{b}$  - known vector, discrete version of  $f(x, y)$

# Sparse linear algebra

Solving

$$A\mathbf{u} = \mathbf{b}$$

- $A$  - sparse matrix, discrete version of spatial influence
- $\mathbf{b}$  - known vector, discrete version of  $f(x, y)$
- $\mathbf{u}$  - unknown vector, discrete version of  $u(x, y)$

# Sparse linear algebra

## Solving

$$Au = b$$

*There are many ways to do this too.  
Iterative methods!*

- $A$  - sparse matrix, discrete version of spatial influence
- $b$  - known vector, discrete version of  $f(x, y)$
- $u$  - unknown vector, discrete version of  $u(x, y)$

# Iterative methods

Rely on matrix-vector multiplications.

- Iterate until a desired tolerance is satisfied

$$|Au_k - b| < tol$$

*tol – user chosen tolerance*

*k – iteration number*

# Iterative methods

Rely on matrix-vector multiplications.

- Iterate until a desired tolerance is satisfied

$$|Au_k - b| < tol$$

Methods:

- Splitting methods

$$A = M - N \rightarrow Mu_{k+1} = b + Nu_k$$

*tol – user chosen tolerance*

*k – iteration number*

# Iterative methods

Rely on matrix-vector multiplications.

- Iterate until a desired tolerance is satisfied

$$|Au_k - b| < tol$$

*tol – user chosen tolerance*

*k – iteration number*

Methods:

- Splitting methods

$$A = M - N \rightarrow Mu_{k+1} = b + Nu_k$$

- Jacobi

[Jacobi1845]

# Iterative methods

Rely on matrix-vector multiplications.

- Iterate until a desired tolerance is satisfied

$$|Au_k - b| < tol$$

*tol – user chosen tolerance*

*k – iteration number*

Methods:

- Splitting methods

$$A = M - N \rightarrow Mu_{k+1} = b + Nu_k$$

- Jacobi [Jacobi1845]
- SOR [Southwell1946]

# Iterative methods

Rely on matrix-vector multiplications.

- Iterate until a desired tolerance is satisfied

$$|Au_k - b| < tol$$

*tol – user chosen tolerance*

*k – iteration number*

Methods:

- Splitting methods

$$A = M - N \rightarrow Mu_{k+1} = b + Nu_k$$

- Jacobi
- SOR

[Jacobi1845]

[Southwell1946]

- Krylov methods

$$u_k \in \text{span}\{b, Ab, A^2b, \dots, A^{k-1}b\}$$

# Iterative methods

Rely on matrix-vector multiplications.

- Iterate until a desired tolerance is satisfied

$$|Au_k - b| < tol$$

*tol – user chosen tolerance*

*k – iteration number*

Methods:

- Splitting methods

$$A = M - N \rightarrow Mu_{k+1} = b + Nu_k$$

- Jacobi [Jacobi1845]
- SOR [Southwell1946]

- Krylov methods

$$u_k \in \text{span}\{b, Ab, A^2b, \dots, A^{k-1}b\}$$

- Conjugate gradient [Hestenes1952]

# Iterative methods

Rely on matrix-vector multiplications.

- Iterate until a desired tolerance is satisfied

$$|Au_k - b| < tol$$

*tol – user chosen tolerance*

*k – iteration number*

Methods:

- Splitting methods

$$A = M - N \rightarrow Mu_{k+1} = b + Nu_k$$

- Jacobi [Jacobi1845]
- SOR [Southwell1946]

- Krylov methods

$$u_k \in \text{span}\{b, Ab, A^2b, \dots, A^{k-1}b\}$$

- Conjugate gradient [Hestenes1952]
- GMRES [Saad1986]

# Pros and cons of iterative methods

Pros:

Cons:

# Pros and cons of iterative methods

Pros:

- Algorithms scale well ( $\mathcal{O}(km)$ , typically with  $k \ll m$ ).

Cons:

# Pros and cons of iterative methods

Pros:

- Algorithms scale well ( $\mathcal{O}(km)$ , typically with  $k \ll m$ ).

Cons:

- Generally can't re-use solutions.

# Pros and cons of iterative methods

## Pros:

- Algorithms scale well ( $\mathcal{O}(km)$ , typically with  $k \ll m$ ).
- Only solve as accurately as you need/want.

## Cons:

- Generally can't re-use solutions.

# Pros and cons of iterative methods

Pros:

- Algorithms scale well ( $\mathcal{O}(km)$ , typically with  $k \ll m$ ).
- Only solve as accurately as you need/want.

Cons:

- Generally can't re-use solutions.
- Spectral properties **may not be known** for a given matrix.

Pro/Con:

- Time to solution dependent on matrix *spectrum*
  - Preconditioning methods are **spectrum-dependent**

# Preconditioning

Multiply linear system

$$Au = b$$

# Preconditioning

Multiply linear system

$$Au = b$$

$$\textcolor{red}{M_L^{-1}Au = M_L^{-1}b}$$

# Preconditioning

Multiply linear system

$$Au = b$$

$$M_L^{-1}Au = M_L^{-1}b$$

$$M_L^{-1}A(M_R M_R^{-1}) u = M_L^{-1}b$$

# Preconditioning

Multiply linear system

$$Au = b$$

$$M_L^{-1}Au = M_L^{-1}b$$

$$M_L^{-1}A\left(M_R M_R^{-1}\right)u = M_L^{-1}b$$

$$\left(M_L^{-1}AM_R\right)\left(M_R^{-1}u\right) = M_L^{-1}b$$

# Preconditioning

Multiply linear system

$$Au = b$$

$$M_L^{-1}Au = M_L^{-1}b$$

$$M_L^{-1}A(M_R M_R^{-1})u = M_L^{-1}b$$

$$(M_L^{-1}AM_R)(M_R^{-1}u) = M_L^{-1}b$$

$$\tilde{A}\tilde{u} = \tilde{b}$$

# Preconditioning

Multiply linear system

$$Au = b$$

$$M_L^{-1}Au = M_L^{-1}b$$

$$M_L^{-1}A(M_R M_R^{-1})u = M_L^{-1}b$$

$$(M_L^{-1}AM_R)(M_R^{-1}u) = M_L^{-1}b$$

$$\tilde{A}\tilde{u} = \tilde{b}$$

# Preconditioning

Multiply linear system

$$Au = b$$

$$M_L^{-1}Au = M_L^{-1}b$$

$$M_L^{-1}A(M_R M_R^{-1})u = M_L^{-1}b$$

$$(M_L^{-1}AM_R)(M_R^{-1}u) = \textcolor{red}{M_L^{-1}b}$$

$$\tilde{A}\tilde{u} = \tilde{b}$$

# Preconditioning

Multiply linear system

$$Au = b$$

$$M_L^{-1}Au = M_L^{-1}b$$

$$M_L^{-1}A(M_R M_R^{-1})u = M_L^{-1}b$$

$$(M_L^{-1}AM_R)(M_R^{-1}u) = M_L^{-1}b$$

$$\tilde{A}\tilde{u} = \tilde{b}$$

Krylov solutions look like

$$\tilde{u}_k \in \text{span}\{\tilde{b}, \tilde{A}\tilde{b}, \dots, \tilde{A}^{k-1}\tilde{b}\}$$

$$u_k = M_R\tilde{u}_k$$

# Preconditioning

Multiply linear system

$$Au = b$$

$$M_L^{-1}Au = M_L^{-1}b$$

$$M_L^{-1}A(M_R M_R^{-1})u = M_L^{-1}b$$

$$(M_L^{-1}AM_R)(M_R^{-1}u) = M_L^{-1}b$$

$$\tilde{A}\tilde{u} = \tilde{b}$$

Krylov solutions look like

$$\tilde{u}_k \in \text{span}\{\tilde{b}, \tilde{A}\tilde{b}, \dots, \tilde{A}^{k-1}\tilde{b}\}$$

$$u_k = M_R\tilde{u}_k$$

... depends on **spectrum of  $\tilde{A}$**  rather than  $A$ .

# Preconditioning

Multiply linear system

$$Au = b$$

$$M_L^{-1}Au = M_L^{-1}b$$

$$M_L^{-1}A(M_R M_R^{-1})u = M_L^{-1}b$$

$$(M_L^{-1}A M_R)(M_R^{-1}u) = M_L^{-1}b$$

$$\tilde{A}\tilde{u} = \tilde{b}$$

Krylov solutions look like

$$\tilde{u}_k \in \text{span}\{\tilde{b}, \tilde{A}\tilde{b}, \dots, \tilde{A}^{k-1}\tilde{b}\}$$

$$u_k = M_R\tilde{u}_k$$

... depends on spectrum of  $\tilde{A}$  rather than  $A$ .

- $M_L^{-1}$  – left preconditioning matrix
- $M_R$  – right preconditioning matrix

Choose so that  $\tilde{A}$  has a *nice* spectrum.

## Specific problem (from PBSM3D)

For stability of time-integration, requires solution to a steady-state PDE at each time step:

$$(\mathbb{I} - D\nabla^2) u = f(x, y)$$

## Specific problem (from PBSM3D)

For stability of time-integration, requires solution to a steady-state PDE at each time step:

$$(\mathbb{I} - D\nabla^2) u = f(x, y)$$

- $D$  – smoothing coefficient
- $u(x, y)$  – rate of snow mass transport at  $x, y$  coordinate
- $f(x, y)$  – known quantity dependent on other physical parameters

## Specific problem (from PBSM3D)

For stability of time-integration, requires solution to a steady-state PDE at each time step:

$$(\mathbb{I} - D\nabla^2) u = f(x, y)$$

*(Requires much more, but we're just going to stick to this again)*

- $D$  – smoothing coefficient
- $u(x, y)$  – rate of snow mass transport at  $x, y$  coordinate
- $f(x, y)$  – known quantity dependent on other physical parameters

# Discretization of specific problem

$$Au = b$$

# Discretization of specific problem

$$Au = b$$

$A$  represents a finite volume discretization of  $(\mathbb{I} - D\nabla^2)$ .  
( $A$  is symmetric positive definite.)

# Discretization of specific problem

$$Au = b$$

$A$  represents a finite volume discretization of  $(\mathbb{I} - D\nabla^2)$ .  
( $A$  is symmetric positive definite.)

Components of  $u$  represent the rate of snow mass transport, constant on each element.

# Discretization of specific problem

$$Au = b$$

$A$  represents a finite volume discretization of  $(\mathbb{I} - D\nabla^2)$ .  
( $A$  is symmetric positive definite.)

Components of  $u$  represent the rate of snow mass transport, constant on each element.

Components of  $b$  represent known quantities dependent on current snow mass (among other things), constant on each element.

## Discretization of specific problem

$$Au = b$$

$A$  represents a finite volume discretization of  $(\mathbb{I} - D\nabla^2)$ .  
( $A$  is symmetric positive definite.)

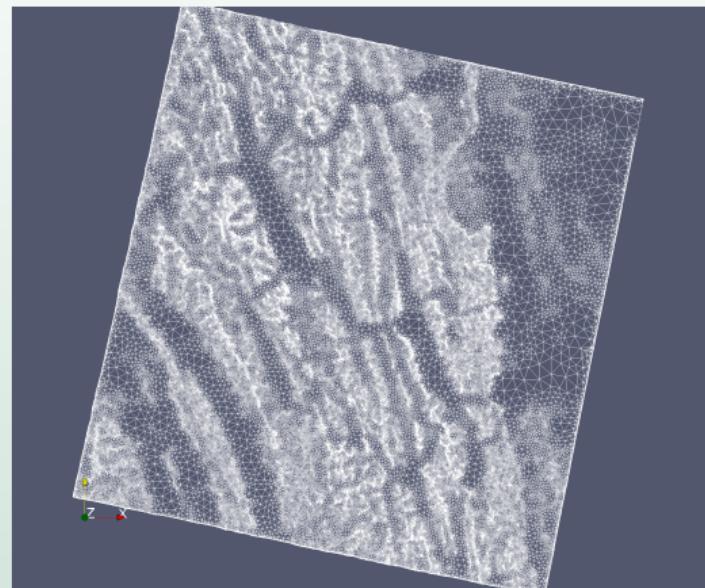
Components of  $u$  represent the rate of snow mass transport, constant on each element.

Components of  $b$  represent known quantities dependent on current snow mass (among other things), constant on each element.

Use known solution  $u^* = 1^{N \times 1}$ ,  $b^* = Au^*$  for testing,  
i.e., solve  $Au = b^*$ .

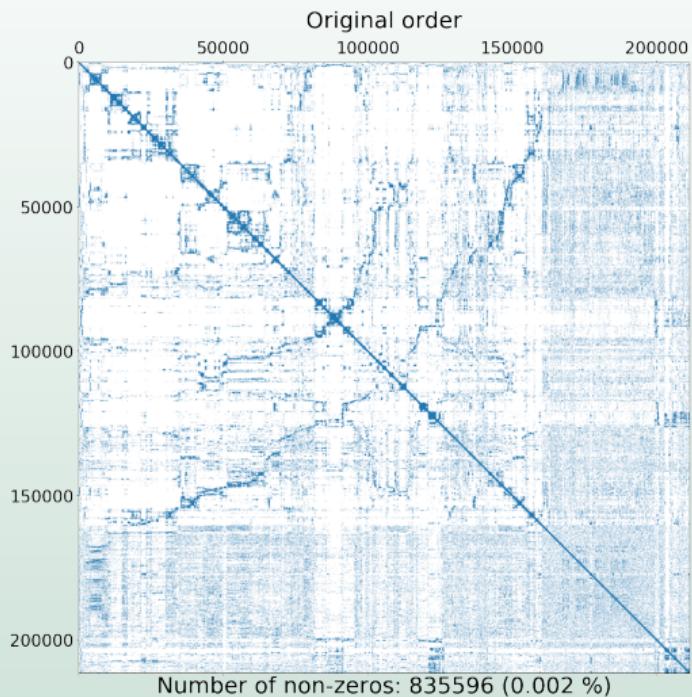
# Mesh: fine resolution

High resolution (211415 elements)



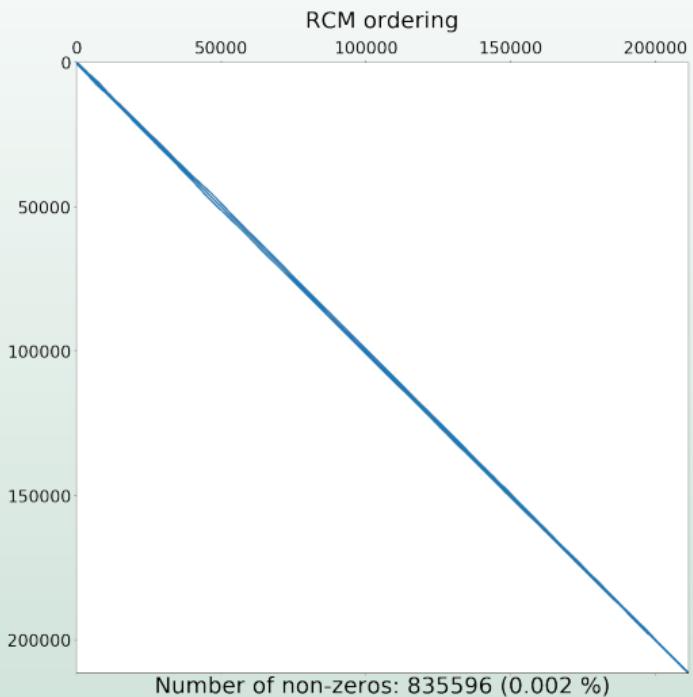
# Matrix structure: original order

Bandwidth: 210855



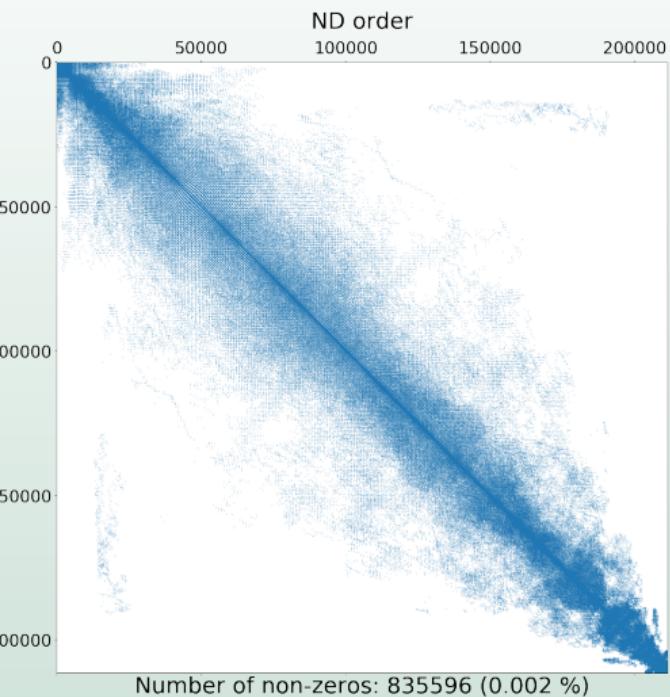
# Matrix structure: RCM order

Bandwidth: 1320



# Matrix: ND order

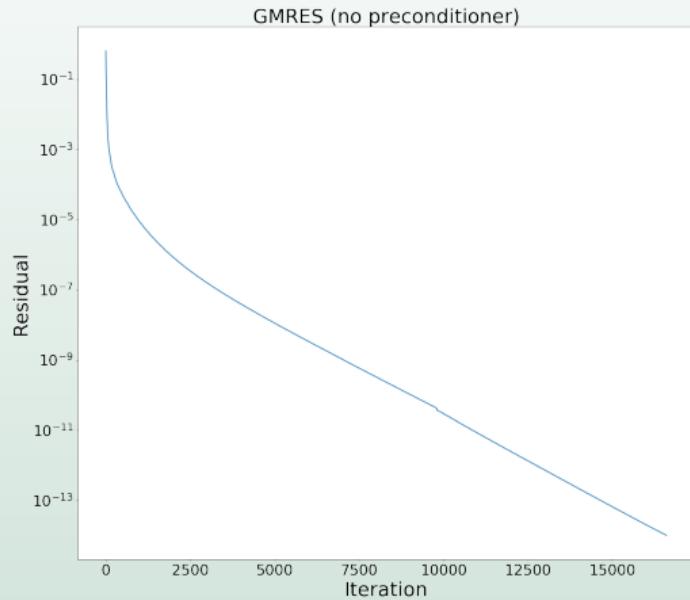
Bandwidth: 173136



# Performance: GMRES

`tol=1e-14, N_restart=100`

No preconditioner.

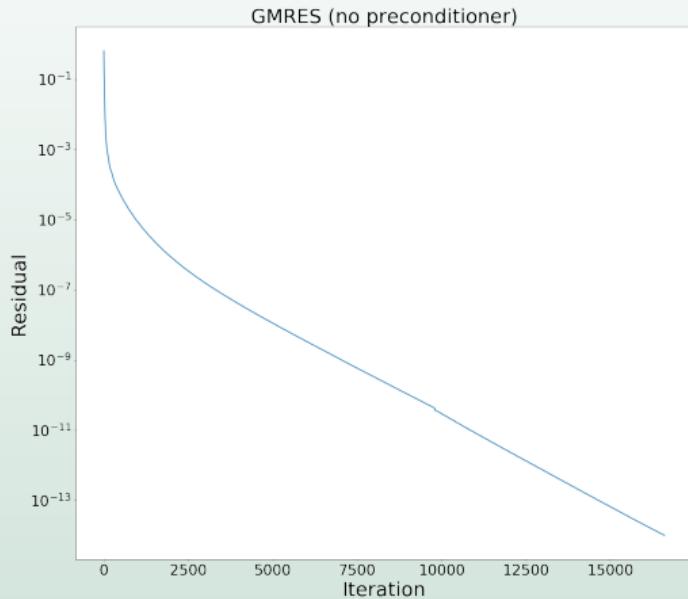


# Performance: GMRES

`tol=1e-14, N_restart=100`

No preconditioner.

Number of iterations: **16607**



# Performance: GMRES

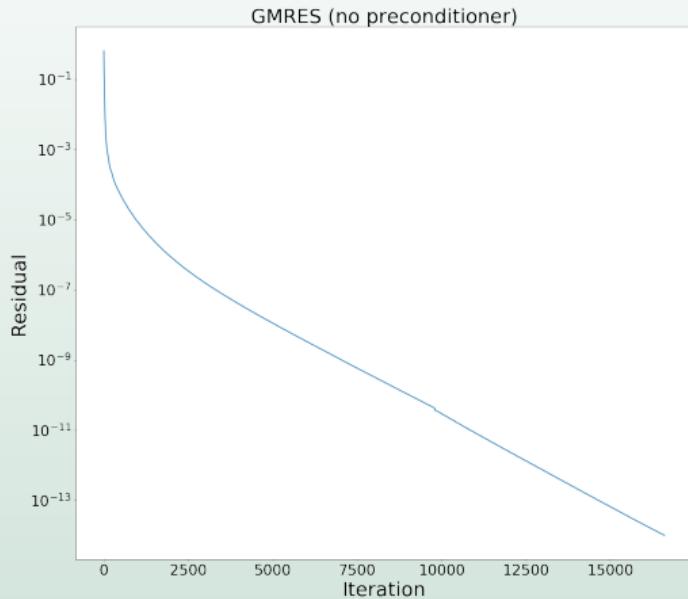
`tol=1e-14, N_restart=100`

No preconditioner.

Number of iterations: 16607

Original order

Wall time: 1min 52s



# Performance: GMRES

`tol=1e-14, N_restart=100`

No preconditioner.

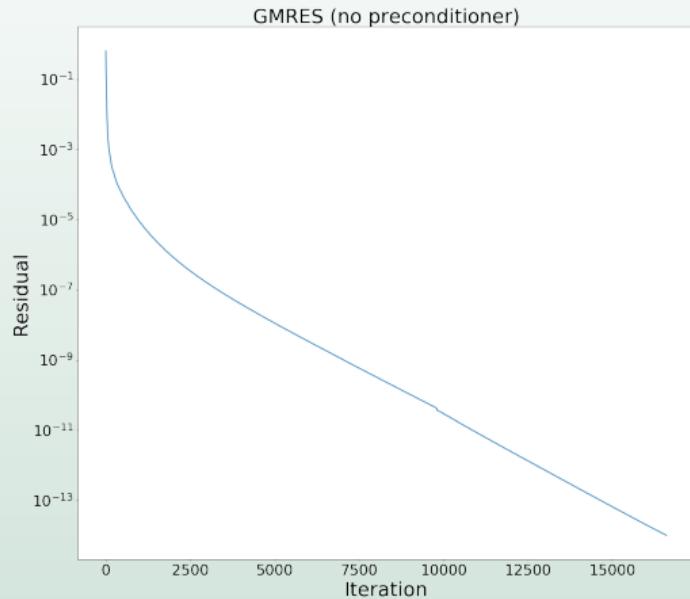
Number of iterations: 16607

Original order

Wall time: 1min 52s

RCM order

Wall time: 1min 23s



# Performance: GMRES

`tol=1e-14, N_restart=100`

No preconditioner.

Number of iterations: 16607

Original order

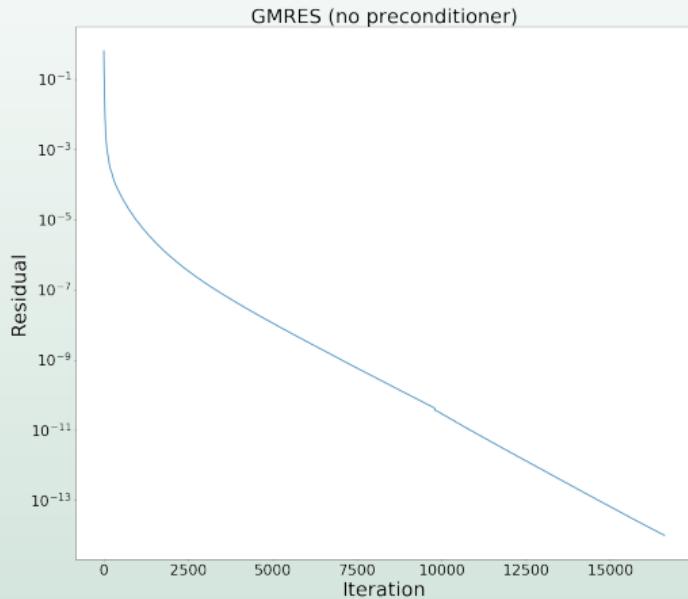
Wall time: 1min 52s

RCM order

Wall time: 1min 23s

ND order

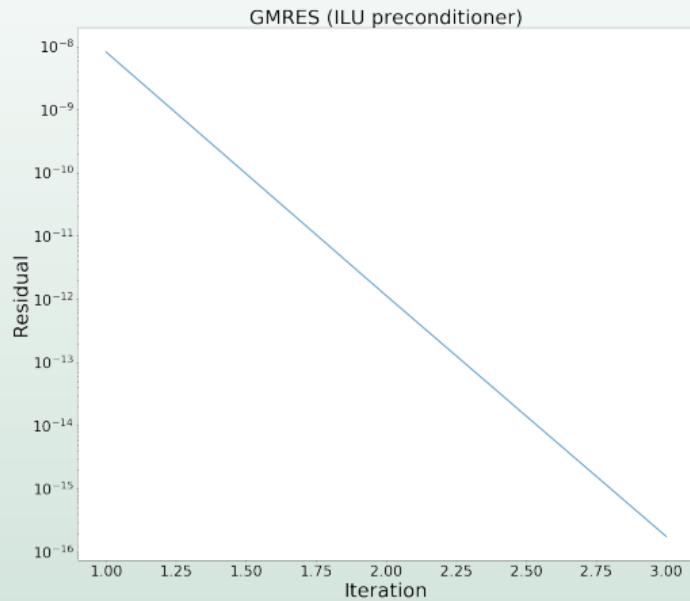
Wall time: 1min 42s



# Performance: GMRES

$\text{tol}=1\text{e-}14$

ILU preconditioner.

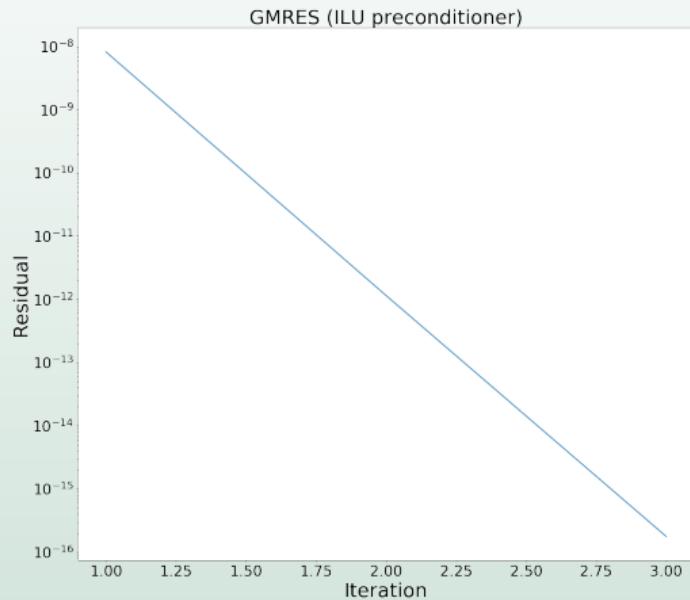


# Performance: GMRES

$\text{tol}=1\text{e-}14$

ILU preconditioner.

Number of iterations: 3



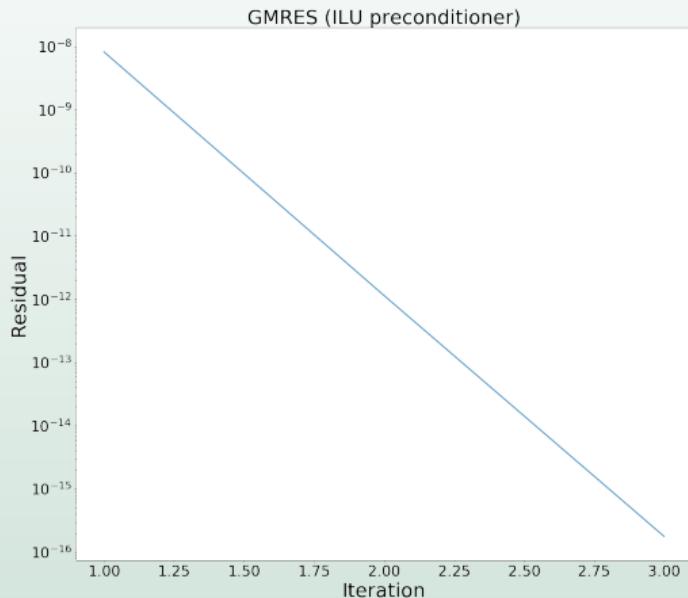
# Performance: GMRES

`tol=1e-14`

ILU preconditioner.

Number of iterations: 3

Order	Setup (s)	Solve (s)	Total (s)
original	0.682	0.497	1.179
RCM	0.602	0.426	1.028
ND	0.644	0.466	1.110



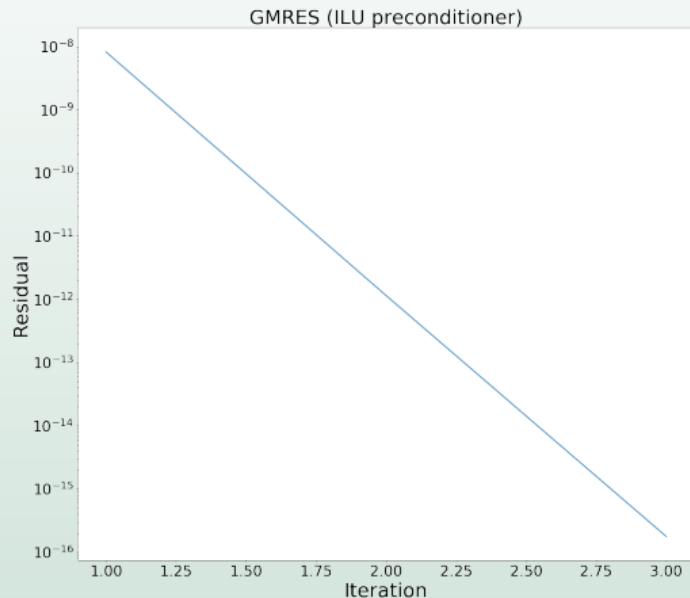
# Performance: GMRES

`tol=1e-14`

ILU preconditioner.

Number of iterations: 3

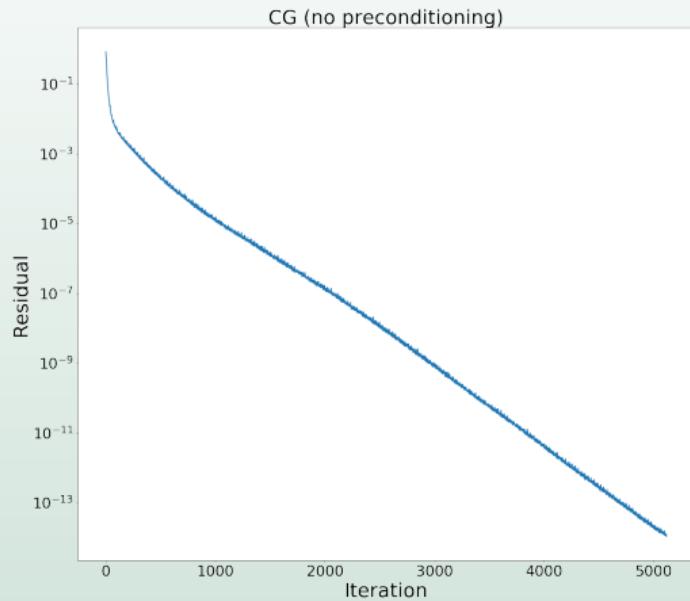
Order	Setup (s)	Solve (s)	Total (s)
original	0.682	0.497	1.179
RCM	0.602	0.426	1.028
ND	0.644	0.466	1.110



# Performance: CG

$\text{tol}=1\text{e-}14$

No preconditioner.

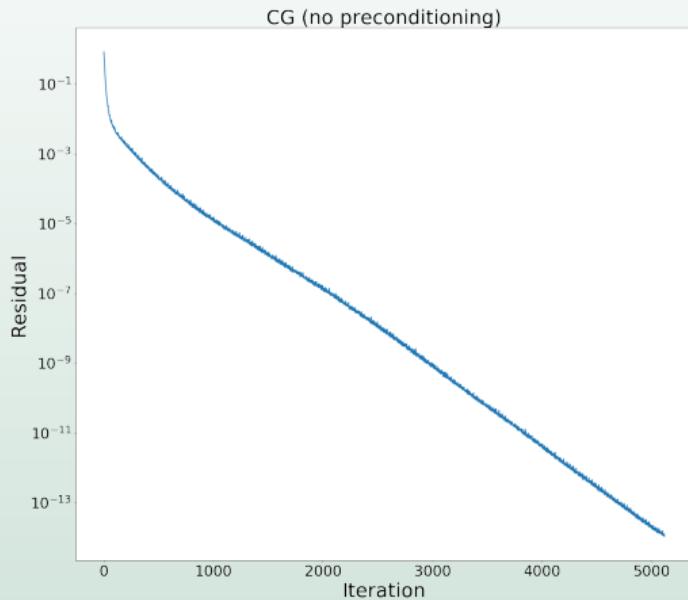


# Performance: CG

$\text{tol}=1\text{e-}14$

No preconditioner.

Number of iterations: **5117**



## Performance: CG

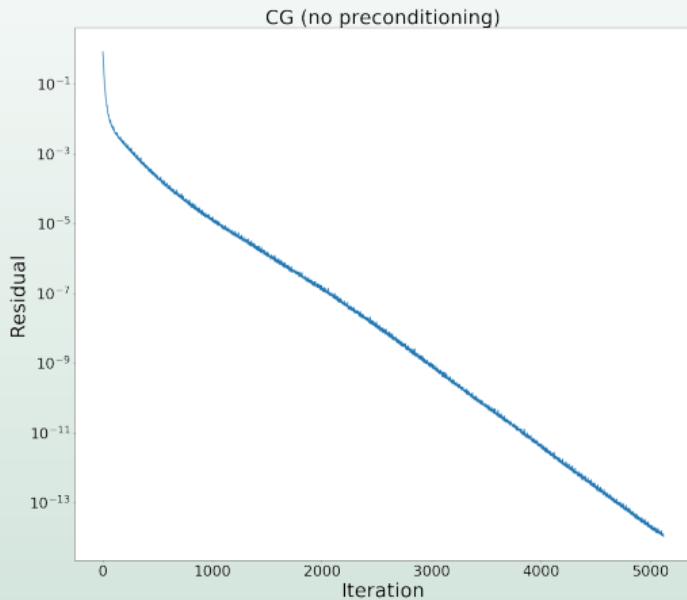
`tol=1e-14`

No preconditioner.

Number of iterations: 5117

Original order

Wall time: **21.3s**



# Performance: CG

$\text{tol}=1\text{e-}14$

No preconditioner.

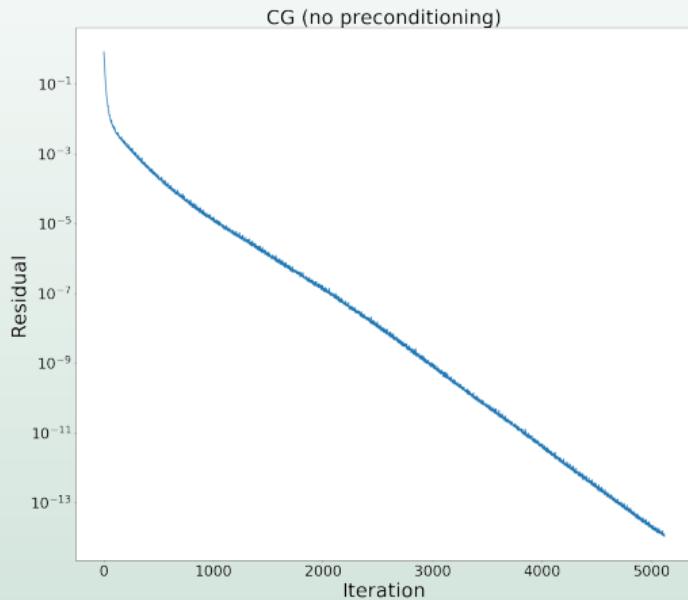
Number of iterations: 5117

Original order

Wall time: 21.3s

RCM order

Wall time: 13.6s



## Performance: CG

`tol=1e-14`

No preconditioner.

Number of iterations: 5117

Original order

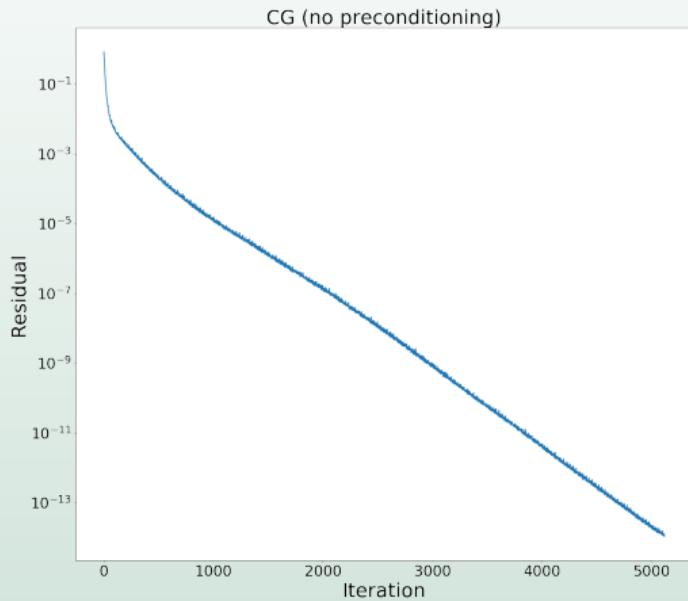
Wall time: 21.3s

RCM order

Wall time: 13.6s

ND order

Wall time: 20.4s



## Performance: CG

`tol=1e-14`

ILU preconditioner.

*No plot*

## Performance: CG

`tol=1e-14`

ILU preconditioner.

Number of iterations: 3

*No plot*

## Performance: CG

`tol=1e-14`

ILU preconditioner.

Number of iterations: 3

Order	Setup (s)	Solve (s)	Total (s)	No plot
original	0.682	0.341	1.023	
RCM	0.602	0.301	0.903	
ND	0.644	0.333	0.977	

## Performance: CG

`tol=1e-14`

ILU preconditioner.

Number of iterations: 3

Order	Setup (s)	Solve (s)	Total (s)	No plot
original	0.682	0.341	1.023	
RCM	0.602	0.301	0.903	
ND	0.644	0.333	0.977	

# Fine mesh summary

211415 elements → 835596 non-zero matrix entries

LU		
Order	Time (s)	Accuracy
original	25145	1.2e-13
RCM	53.3	1.6e-13
ND	1.3	1.5e-13

GMRES		
	Time (s)	Accuracy
	1.179	6.3e-10
	1.028	5.9e-10
	1.110	6.1e-10

CG		
	Time (s)	Accuracy
	1.023	1.2e-13
	0.903	1.2e-13
	0.977	1.1e-13

# Fine mesh summary

211415 elements → 835596 non-zero matrix entries

LU		
Order	Time (s)	Accuracy
original	25145	1.2e-13
RCM	53.3	1.6e-13
ND	1.3	1.5e-13

GMRES		
	Time (s)	Accuracy
	1.179	6.3e-10
	1.028	5.9e-10
	1.110	6.1e-10

CG		
	Time (s)	Accuracy
	1.023	1.2e-13
	0.903	1.2e-13
	0.977	1.1e-13

Overall Winner: Preconditioned CG

# Fine mesh summary

211415 elements → 835596 non-zero matrix entries

LU		
Order	Time (s)	Accuracy
original	25145	1.2e-13
RCM	53.3	1.6e-13
ND	1.3	1.5e-13

GMRES		
	Time (s)	Accuracy
	1.179	6.3e-10
	1.028	5.9e-10
	1.110	6.1e-10

CG		
	Time (s)	Accuracy
	1.023	1.2e-13
	0.903	1.2e-13
	0.977	1.1e-13

Overall Winner: Preconditioned CG

Notable showing: LU with ND order

*If many solves can be performed from a single factorization.*

# Exercise

You have the matrix in a text file in

<https://github.com/uofs-simlab/Math313>.

Load that into something like PETSc/SLEPc to play around with solving algorithms and eigenvalue computations.

# Concluding remarks

- A lot of pieces have to come together in the real world before we get to linear algebra.

# Concluding remarks

- A lot of pieces have to come together in the real world before we get to linear algebra.
- There's still a lot to investigate once we get there.

# Concluding remarks

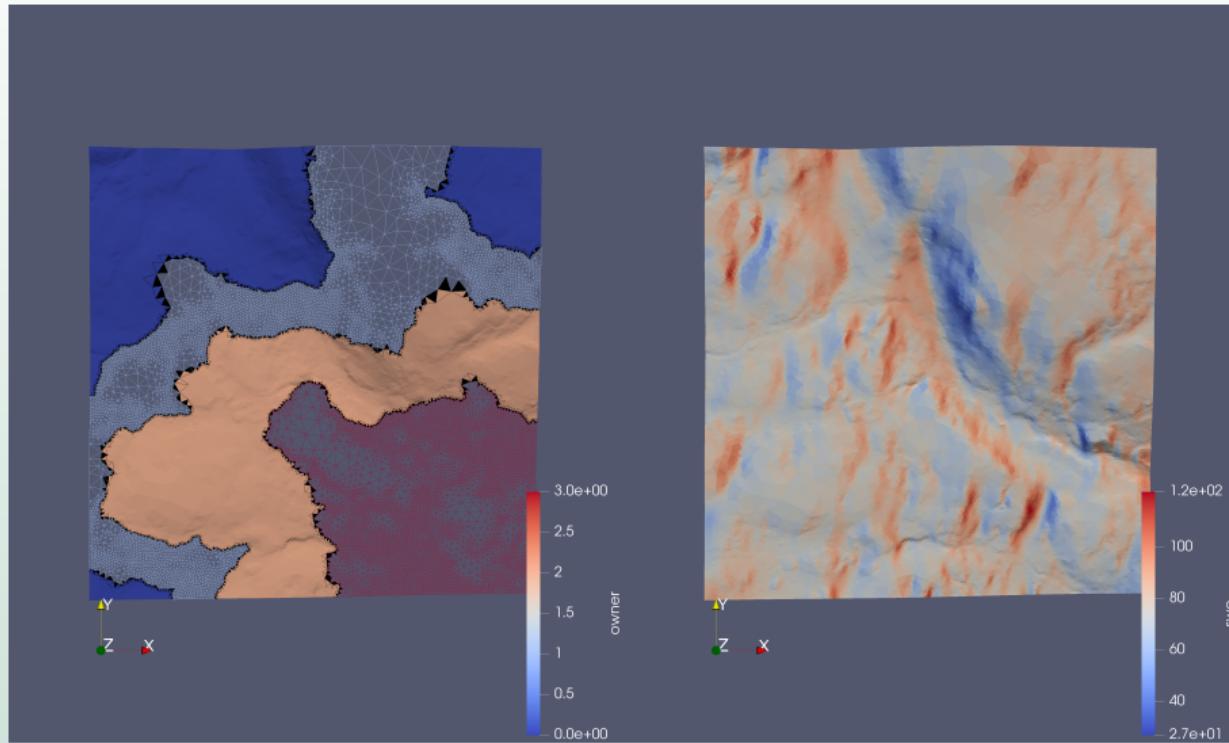
- A lot of pieces have to come together in the real world before we get to linear algebra.
- There's still a lot to investigate once we get there.
  
- Substantial performance improvement obtained through preconditioning.

# Distributed linear algebra

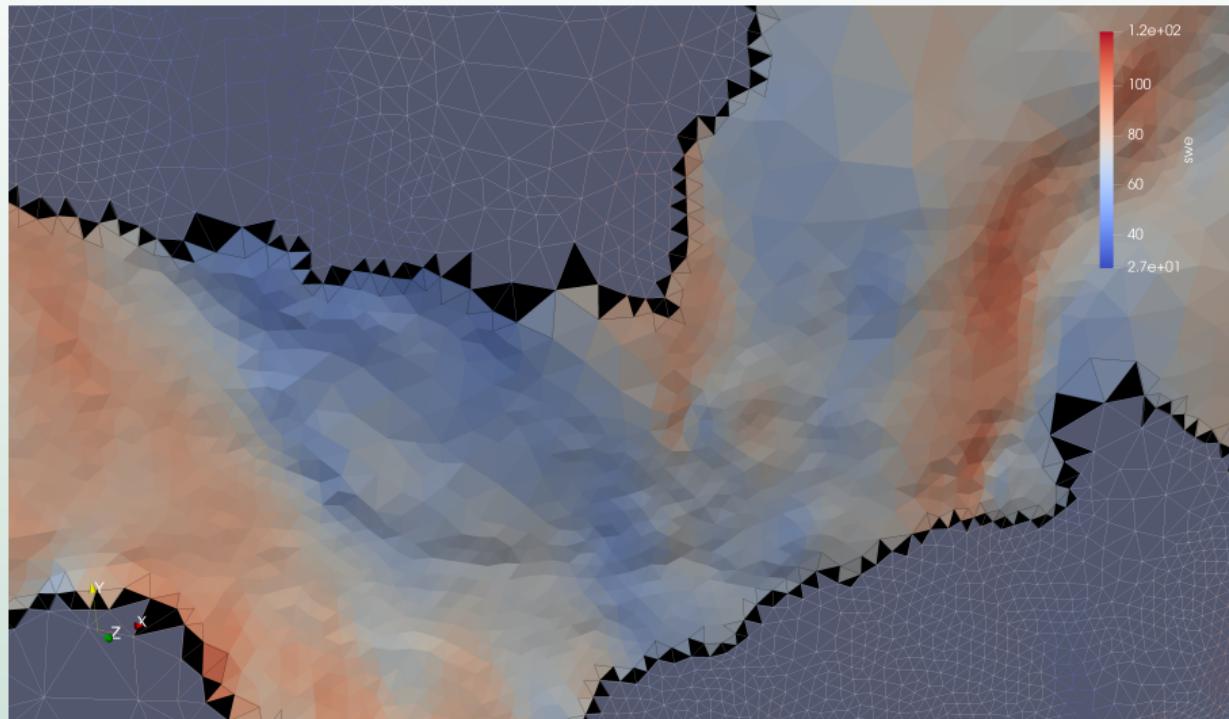
*Done within CHM, not in the jupyter notebook.*

<https://github.com/ChrisMarsh/CHM/tree/dev>

# Ownership of distributed mesh

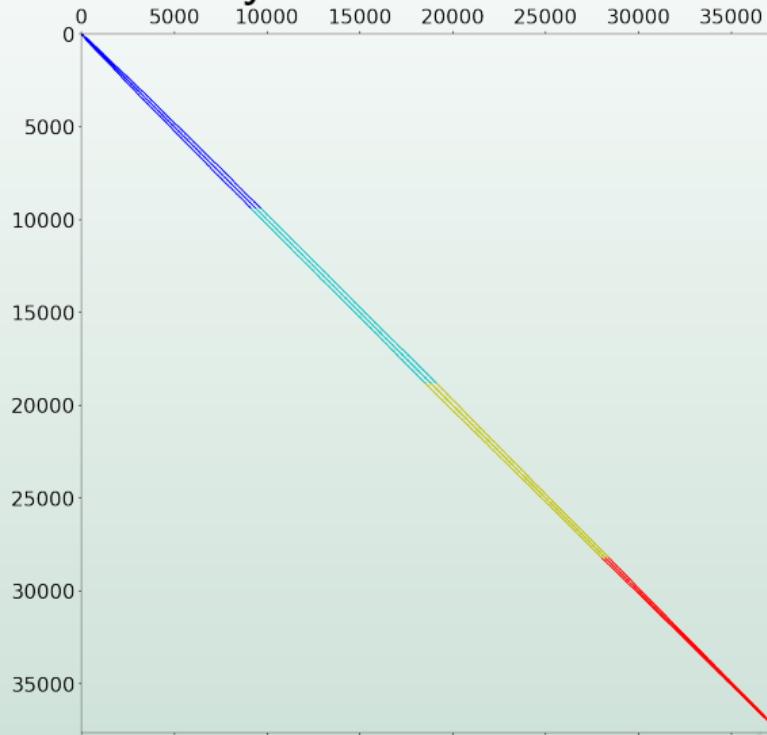


... introduces ghosted elements



# ... and distributed Linear Algebra

## Distributed system



# Performance

## Linear system iterations

Some preconditioners aren't as effective as the number of distributed processes are increased.

Np	Iterations*
1	11
2	34
4	35
8	35
16	42

\* Trilinos/Belos: GMRES(1e-8)  
Trilinos/Itpack2: BlockJacobi(Np)\*ILUT(3,1e-4)